

Automated Sharded MongoDB Deployment and Benchmarking for Big Data Analysis

Mark McCombe, Gregor von Laszewski, Geoffrey C. Fox
Indiana University, Smith Research Center, 2805 E 10th St, Bloomington, Indiana, 47408
laszewski@gmail.com

ABSTRACT

Project CH-818664, KVM: Using Python, Ansible, Bash Shell, and Cloudmesh Client a fully automated process is created for deploying a configurable MongoDB sharded cluster on Chameleon, FutureSystems, and Jetstream cloud computing environments. A user runs a single Python program which configures and deploys the environment based on parameters specified for numbers of Config Server Replicas, Mongos Instances, Shards, and Shard Replication. The process installs either MongoDB version 3.4 or 3.2 as requested by the user. Additionally, functionality exists to run benchmarking tests for each deployment, capturing statistics in a file as input for python visualization programs, the results of which are displayed in this report. These reports depict the impact of MongoDB version and degrees of sharding and replication on performance. Key performance findings regarding version, sharding, and replication are abstracted from this analysis. As background, technologies and concepts key to the deployment and benchmarking, such as MongoDB, Python, Ansible, Cloudmesh Client, and Openstack are examined while comparing and using them within different clouds.

KEYWORDS

Chameleon CCloud, Jetstream, Futuresystems, MongoDB, Cloud Computing, Ansible, Python, Cloudmesh Client, Openstack

1 INTRODUCTION

Three clouds were selected for deployment: Chameleon Cloud, Futuresystems (also referred to as Kilo in some sections of this document), and Jetstream. In our automated deployment and benchmarking process, the cloud name is passed as a parameter to the deploy function of the main script and a customized version of MongoDB is deployed to the selected cloud.

2 CLOUD USAGE

We compared within the allocation limitations of a class multiple cloud performances by varying a number of parameters. In addition to Chameleon cloud we also used Jetstream and the Futuresystems cloud. Adding these clouds was essential to obtain comparisons of chameleon cloud to other clouds. A full report is available as part of the Class proceedings.

Table 1 shows a comparison of key server computing resources on Chameleon, FutureSystems, and Jetstream cloud environments.

3 PROJECTS

4 EXPERIMENTAL CONFIGURATION

Cloudmesh client is used to simplify management of vms across different clouds. The Cloudmesh Client toolkit is an open source

Table 1: Cloud Server Hardware Specification Comparison
[1] [3] [2]

	FutureSystems	Chameleon	Jetstream
CPU	Xeon E5-2670	Xeon X5550	Haswell E-2680
cores	1024	1008	7680
speed	2.66GHz	2.3GHz	2.5GHz
RAM	3072GB	5376GB	40TBr
storage	335TB	2TB	2 TB

client interface that standardizes access to various clouds, clusters, and workstations [4]. Cloudmesh Client is a python based application. In the deployment, Cloudmesh Client is used to handle most interaction with the Virtual Machines in the clouds. Cloudmesh Client provides functionality in three main areas: Key Management, OpenStack Security, and virtual machine management. For key management, Cloudmesh's key add and upload commands simplify secure interaction with the cloud environments. For Openstack security, Cloudmesh's secgroup commands allow new security rules to be added and uploaded to the cloud. Virtual machine management is performed with Cloudmesh's cluster functionality, which allows easy creation and deletion of virtual machines and communication between them. Cloudmesh Client simplifies and standardized interaction with the cloud for these tasks. This allows us to more easily port the deployment to additional clouds that are supported by Cloudmesh. Furthermore, by encapsulating the logic necessary to perform these tasks we are shielded from changes in interfaces made by individual clouds.

4.1 Resource Requirements

As this project was a class project the available VM hours were limited. However, we have been able to conduct a significant comparison given the restrictions.

4.2 Capability Requirements

The main feature we needed for this project was the creation of VMs and the execution of our applications within these VMs

4.3 Monitoring Requirements

Monitoring and benchmarking was conducted by hand without need for specialized services.

4.4 Features offered by Chameleon Cloud

Chameleon provided one of three clouds to the project.

4.5 New software created

As part of this class we improved the cloudmesh client software [5][6] [7] that was essential to the success of the class.

4.6 Performance Comparison

We have conducted a significant performance comparison among all clouds. However, in this document we only list a few highlights.

4.7 Computing Resources

In all cases, virtual machines are deployed with the Ubuntu 16.04 LTS (Xenial Xerus) operating system. On Openstack the flavor or the machine determines the amount of computing resources (CPU, memory, storage) allocated to it. In our testing, m1.medium was used as the flavor for Chameleon Cloud and FutureSystems, while m1.small was used on Jetstream. Jetstream has more resources allocated to each flavor than Chameleon and FutureSystems, which are similar. In order to perform similar tests on each cloud, flavors with identical CPU and memory were selected. Table ?? shows the comparative resources of the flavors used in our testing. While storage is lower on Jetstream, it is sufficient for our tests and should not significantly impact performance.

5 DEPLOYMENT EXAMPLES

The configuration parameters and cluster and Ansible deployment times are captured in a file for each deployment (benchmarking timings are later captured as well).

Deployment A constitutes a simple deployment with only one of each component being created. This deployment may only be suitable for a development or test environment. Deployment A completed in 330 seconds.

Deployment B constitutes a more complex deployment with production like replication factors for Config Servers and Shards and an additional Mongos instance. This deployment may be suitable for a production environment as it has greater fault tolerance and redundancy. Deployment B took 1059 seconds to deploy.

Deployment C focused on high performance. It has a high number of shards, nine, but no fault tolerance or redundancy. The deployment may be suitable where performance needs are high and availability is less critical. Deployment C finished in 719 seconds.

5.1 Benchmarking Analysis

5.1.1 Cloud Analysis. Chameleon Cloud was significantly more stable and reliable than FutureSystems and Jetstream Clouds for our testing. Chameleon yields the fastest and most consistent results with very few errors. Jetstream initially had stability problems that were eventually resolved by the Jetstream support team. Once these issues were resolved, Jetstream performance and stability was very close to Chameleon's. FutureSystem performance was the poorest with respect to run time. Environmental errors were initially frequent, but after allocating new floating IPs test would be completed successfully. JetStream performance was good, but the environment was very unstable. Due to its stability and performance, Chameleon was chosen as the environment to test MongoDB version 3.4 versus 3.2, due to its stability.

5.1.2 Impact of Sharding on Reads. Figure 1 depicts the impact on performance of various numbers of shards on a find command in Chameleon, FutureSystems, and Jetstream Clouds. All three clouds show a strong overall decline in run time as the number of shards increases, which shows the positive impact of sharding on performance. For all clouds, reads were over 35 seconds for one shard and less than 10 seconds for five shards. This is a significant gain in performance.

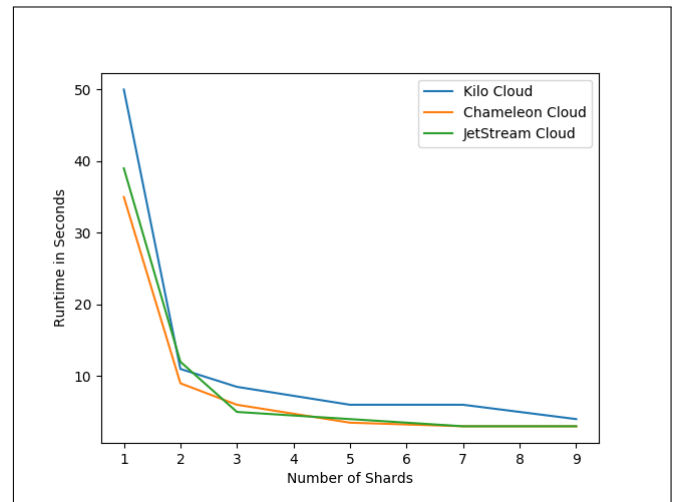


Figure 1: Find Command - Sharding Test

All three clouds show a particularly large gain in performance when increasing from one shard to two. Run time for two shards is less than one third the run time of one shard. Increases in shards beyond two show much smaller incremental gains.

Performance on Chameleon Cloud and Jetstream is very similar for the find test. Kilo performance is worse, although proportionately better than on the mongoimport test. This is an interesting observation as for both deployment and mongoimport, performance was much better on Chameleon and Futuresystems than Kilo. One difference from the mongoimport test is that much less data is being sent over the network. Network speeds could be a factor in this discrepancy.

5.1.3 Impact of Sharding on Writes. Figure 2 depicts the impact on performance of various numbers of shards on a mongoimport command in the three clouds. For all clouds, run time of the mongoimport command in our tests does not appear to be impacted by the number of shards. Since the same amount of data is written with more computing resources available when there are more shards, we might expect to see a performance gain. However, there are possible explanations for performance not improving. First, the mongoimport command may not write data in parallel. This is not indicated in the documentation, but it seems likely that it reads the file serially. Second, resources on the server the data is written to may not be the bottleneck in the write process. Other resources like the network time seem more likely to be the bottleneck. Since we are always going over the network from the mongos instance

to a data shard, regardless of the number of shards, a bottleneck in the network would impact all shard configurations equally.




Figure 2: Mongoimport Command - Sharding Test

While sharding did not benefit a single threaded mongoimport command, it is likely it would benefit other heavy write operations, particularly coming through multiple mongos instances. In a non-sharded environment, this would lead to a heavy load on the single data shard. In a sharded environment, the load on each shard would drop as the number of shards increased.

While performance on Chameleon and FutureSystems was very similar for the find command, performance of the mongoimport command was significantly better on Chameleon than on Kilo. We see approximately 50% better performance on both Chameleon and Jetstream Clouds compared to FutureSystems. Jetstream performance is slightly better than Chameleon for the import test.

5.1.4 Impact of Sharding on MapReduce. Figure 3 shows the performance of MapReduce across various sharding configurations on our three clouds. These results are relatively similar to the find results. While results are inconsistent, particularly on Futuresystems, likely due to environmental issues, all clouds show an overall decrease in processing time with addition of shards. Relative to Mongoimport performance, performance is more similar across the three clouds for MapReduce.

5.1.5 Impact of Replication on Reads. Figure 4 depicts the impact on performance of various numbers of replicas on a find command in Chameleon, FutureSystems, and Jetstream Clouds. These results show no correlation between the number of replicas and find performance.

Similarly to other tests, performance on Chameleon was best for the majority of the test runs in the find replication test, followed by Jetstream, with Futuresystems performing the worst.

5.1.6 Impact of Replication on Writes. Figure 5 depicts the impact on performance of various numbers of replicas on a mongoimport command on our three Clouds. The results show poorer write performance as the number of replicas increase. Given that an extra




Figure 3: MapReduce - Sharding Test




Figure 4: Find Command - Replication Test

copy of data is written with each increase in the replication factor, this performance hit is expected.

Performance on Jetstream and Chameleon were very close on this test with Chameleon only performing significantly better with four or more replicas. FutureSystems import performance was by far the worst of the three clouds.

5.1.7 Impact of Replication on MapReduce. As shown in Figure 6, replication appears to have no impact on MapReduce operations. While there are variations in FutureSystems and Jetstream performance for different numbers of replicas, they do not follow a consistent pattern and appear to be caused by environmental issues. This is an interesting result as increased levels of replication came with a performance penalty for the find command, which also reads data.




Figure 5: Mongoimport Command - Replication Test




Figure 7: Find Command - Version 3.2 vs 3.4




Figure 6: MapReduce - Replication Test




Figure 8: Mongoimport Command - Version 3.2 vs 3.4

As with several other tests, Chameleon MapReduce performance was the best, followed by Jetstream, with FutureSystems again being the worst.

5.1.8 Impact of Version and Sharding on Reads. Figure 7 shows the MongoDB version 3.4 and 3.2 find performance on Chameleon Cloud. Results are very close, with version 3.2 having the best performance for one shard and performance being similar for all other sharding levels.

5.1.9 Impact of Version and Sharding on Writes. Figure 8 shows the MongoDB version 3.4 and 3.2 Mongoimport performance on Chameleon Cloud. Runtimes are similar for each version. Version 3.2 is slightly faster at the lowest sharding levels and Version 3.4 is slightly faster at the highest sharding level. Given the mixed results and close run times, neither version shows a significant advantage for write operations.

5.1.10 Impact of Version and Sharding on MapReduce. Figure 9 shows the MongoDB version 3.4 and 3.2 Mongoimport performance on Chameleon Cloud. Runtimes are similar for each version and with each version being faster at some shard level, which appears to be random. Given the mixed results and close run times, neither version shows a significant advantage for MapReduce operations.

6 CONCLUSION

We have created, tested, and demonstrated a fully automated program to configure and deploy a sharded MongoDB cluster to three cloud environments: Chameleon, Jetstream, and FutureSystems. Using a combination of Python, Bash, and Cloudmesh Client, the a cluster is dynamically deployed with a selected number of Config Server Replicas, Mongos Routers, Shards, and Shard Replicas and either MongoDB version 3.4 or 3.2. Functions also exist for terminating the environment, reporting on data distribution, benchmarking,




Figure 9: MapReduce - Version 3.2 vs 3.4

and reporting on performance testing. An automated benchmarking process to show the impact of well distributed data across shards of a large data set has been run for various configurations. The impact of MongoDB version 3.4 versus 3.2, Sharding, and Replication on performance have been assessed. Testing showed performance and stability on Chameleon Cloud to be the best of our three cloud environments with Jetstream a close second after an operational issue was resolved by the support team. Futuresystems performance consistently lagged behind the other two clouds due to its older hardware. A key finding is that read performance, typically a high priority for noSQL data stores and Big Data operations, increases significantly as shards are added. Testing also showed that a predictable performance penalty is associated with replication. Our comparison of version 3.4 and 3.2 showed no significant differences between version 3.2 and 3.4 performance across various sharding levels.

REFERENCES

- [1] Chameleon. 2016. Hardware Description. web page. (Nov. 2016). <https://www.chameleoncloud.org/about/hardware-description/>
- [2] JetStream. 2016. System Specs. web page. (April 2016). <http://www.jetstream-cloud.org/leadership.php>
- [3] Gregor von Laszewski. 2013. Hardware. web page. (Jan. 2013). <http://futuregrid.github.io/manual/hardware.html>
- [4] Gregor von Laszewski. 2016. Cloudmesh Client Toolkit. web page. (Sept. 2016). <http://cloudmesh.github.io/client/>
- [5] Gregor von Laszewski. 2017. Cloudmesh Client Framework. Github. (2017). <https://github.com/cloudmesh/client>
- [6] Gregor von Laszewski. 2017. Cloudmesh CMD5. Github. (2017). <https://github.com/cloudmesh/cloudmesh.cmd5>
- [7] Gregor von Laszewski. 2017. Cloudmesh REST Framework. Github. (2017). <https://github.com/cloudmesh/cloudmesh.rest>