

RabbitMQ

ABHISHEK GUPTA^{1,*}

¹ School of Informatics and Computing, Bloomington, IN 47408, U.S.A.

* Corresponding authors: abhigupt@iu.edu

project-001, September 21, 2017

RabbitMQ provides messaging platform which allows applications pass messages in a reliable and fault tolerant way. RabbitMQ implements Advanced Message Queuing Protocol (AMQP) and is written in Erlang programming language. It runs on all major operating systems and supports SDK in all major programming languages[1] including objective-C, swift and node.js. When we look for messaging, we look for certain features: asynchronous messaging, large scale, reliability, clustering, multi-protocol, highly available, fault tolerant. RabbitMQ[2] fulfills these requirements and provides a distributed, persistent, highly available, fault tolerant messaging system which can scale as data grows. © 2017

<https://creativecommons.org/licenses/>. The authors verify that the text is not plagiarized.

Keywords: Cloud, I524

<https://github.com/cloudmesh/sp17-i524/blob/master/paper1/S17-IO-3005/report.pdf>

1. INTRODUCTION

RabbitMQ [2] based off AMQP (Advanced Message Queuing Protocol[3]) Which defines how client applications connect to message brokers. Message brokers receive messages from producers and make it available to consumers. The producer posts messages to exchanges and broker agent read these messages from the exchange and write to queues. Consumers can further read messages from the queue. It also supports a notion of acknowledgments where consumers can post an acknowledgment back to broker and which in turn can post messages back to the producer. Broker plays an important role, it eliminates the well-known producer-consumer problem [4]. It also provides persistence in case of failures.

Looking back at the history [5] of the development of RabbitMQ, it started with IBM Miseries in 1993, 1997 Microsoft MQ, 2001 java messaging service. These technologies were incubated in different companies with a common goal to create a messaging service for a distributed system. It was then in 2003 where JP Morgan created the first version AMQP which became the base for RabbitMQ technologies formed in 2006.

2. ARCHITECTURE

At high level, producers publish messages to the broker. Consumers consume messages from the broker. Broker acts like a middle man, who has knowledge of all available exchanges and queues [5]. It maintains bindings between exchanges and queues. Exchanges here are often compared to post offices or mailboxes. Following table shows different type of exchanges supported by RabbitMQ and corresponding default exchanges [2].

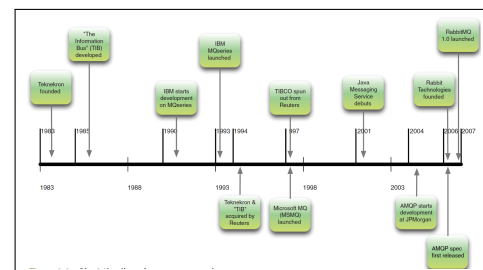


Fig. 1. Timeline for evolution of RabbitMQ [5]

An exchange is a binding to queue(s). So when you send a message to an exchange, the broker checks the routing key of the message and compares it against the queue bindings. Further it sends the message to appropriate queue.

Name	Default pre-declared names
Direct exchange	(Empty string) and amq.direct
Fanout exchange	amq.fanout
Topic exchange	amq.topic
Headers exchange	amq.match (and amq.headers in RabbitMQ)

3. TYPE OF EXCHANGES

The sections below explains various types of exchanges supported by RabbitMQ.

3.1. Default Exchange

It is created by the broker and all queues are bound to default exchange unless specified separately. For example we have a queue called demoqueue, all messages assigned to demoqueue will be routed by default exchange.

3.2. Direct exchange

It is used to route messages to queue with a given routing key. For example a message queue has routing key K. A message with same routing key K will be delivered to same message queue.

3.3. Fanout exchange

Delivers messages to all queues bound to a given exchange. Fanout exchange ignores the routing key. For example, if there are 5 queues bound to an exchange 'E'. Messages to exchange 'E' are delivered to all 5 queues.

3.4. Topic exchange

It delivers messages to a given queue, not just based on the routing key but a pattern specified in the message. This pattern is called topic. It can be useful in scenarios where consumer decides which topic(s) it is interested in. Further, it can subscribe to those topic(s).

3.5. Header exchange

This exchange ignores the routing key but rather uses a header parameter to decide which message goes to which queue. A message is matched when a value in header matches with the one specified in the queue binding. Exchanges have other important attributes apart from routing key and exchange type. For example, name, durability, auto-delete, arguments etc. Durability allows messages to persist on disk in case broker restarts. Auto-delete deletes the exchange once all queues have finished using the exchange.

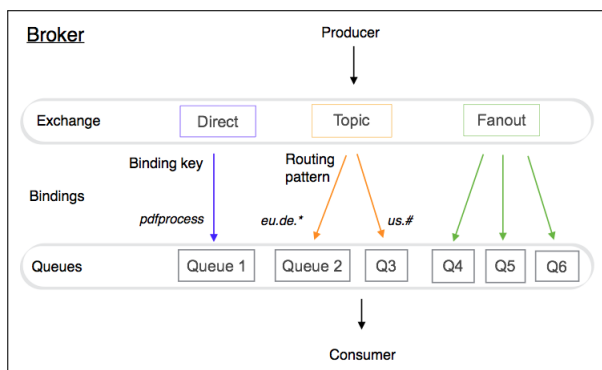


Fig. 2. Different type of exchanges supported by RabbitMQ[6]

4. BINDINGS

Binding is a rule that defines how exchanges will route messages to the queue. To route all messages from an exchange E to queue Q, Q has to be bound to Exchange E. Binding uses routing key as one of the criteria to route messages to a queue. However, routing key is optional and not always applicable.

5. CONSUMERS

Messages from message queue are eventually used or consumed by consumers. Consumers can use push or pull mechanism to consume these messages. Push API have messages delivered to the consumer whereas a Pull API is used to fetch messages from the queue.

6. MESSAGE ACKNOWLEDGEMENT

RabbitMQ has a built-in mechanism to send and receive acknowledgements. Producer can send a message and wait for the acknowledgement in the response queue of the consumer. Once consumer successfully receives the message, it can post an acknowledgement to the response queue. Here request queue and request exchange can be used to send and receive messages. Response queue and response exchange can be used to send and receive acknowledgements. This mechanism makes RabbitMQ robust to handle failure scenarios.

7. CLUSTERING

To achieve high availability [5] and making sure producers and consumers send and receive data without knowing about node failures, clustering was introduced. RabbitMQ follows OTP(open telecom platform) framework provided by erlang to achieve high availability. RabbitMQ by default doesn't replicate the content of queues i.e. all queues are stored on one node in the cluster. To achieve clustering RabbitMQ keeps track of metadata for queue, exchange, binding and vhost. In case of cluster mode, where broker, producer and consumer run on different nodes, RabbitMQ only stores information about the queue like metadata, state and contents on one node rather than all nodes on the cluster. However, it stores metadata and pointer to actual data on each node on the cluster. This is to optimize storage space and performance.

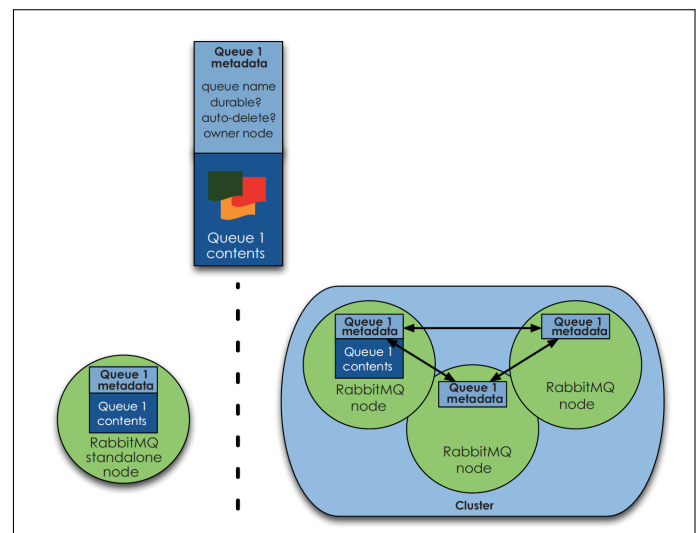


Fig. 3. Shows queue metadata for a cluster vs single node [5]

Since, exchanges act like lookup tables for queue bindings, they are replicated across all nodes on the cluster.

8. MANAGEMENT

RabbitMQ provides all management using:

- Web UI
- REST interface
- rabbitmqctl command line utility

Web UI can be used by the administrator to create user, monitor queues and exchanges, view statistics, add configurations etc. Similar functionality can be achieved by CLI(command line interface) utility called rabbitmqctl or using REST API. Rabbitmqctl can be used to automate RabbitMQ deployments and management. It can also be used to write automated tests. REST API can be used for integrating with 3rd part UI and plugins. Using REST API, you can monitor the number of connections, download or upload a configuration, list the nodes in the cluster, create or delete RabbitMQ users, view or create virtual host, set permission for a user etc. You can also list all current APIs using '/api' URI.

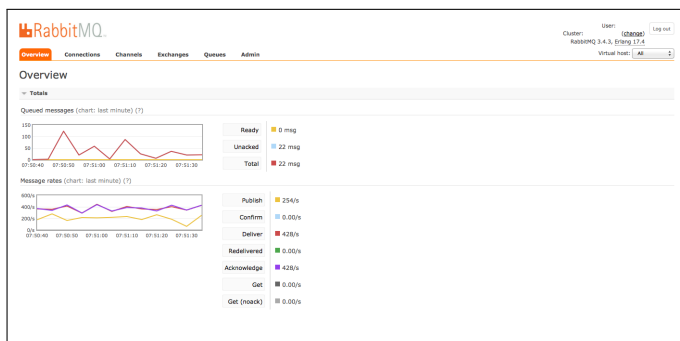


Fig. 4. RabbitMQ management UI showing messages queued and message data rates[6]

9. LICENSING

RabbitMQ is licensed under Mozilla Public License(MPL) and GPL v2 [7]. MPL license [8] allows "for free use, modification, distribution, and "exploit[ation]" of the work, but does not grant the licensee any rights to a contributor's trademarks".

10. USE CASES

RabbitMQ messaging can be useful for applications which require asynchronous messaging i.e. an application initiates a task by posting a message to RabbitMQ, it does not have to wait for the task to get completed. Application can periodically check for the status of completion of the task. RabbitMQ can scale up as the data needs of the application grow, by just adding additional nodes to the RabbitMQ cluster. When applications run as micro-services in a containerized environment, RabbitMQ is useful to communicate and share data between different services. RabbitMQ can be managed separately using management Web UI, CLI tool and REST API. This decouples the messaging layer from application and makes the overall design robust.

11. CONCLUSION

RabbitMQ is an open source platform and provides a robust messaging platform for applications. It provides simple manageability decoupling with the application. RabbitMQ can scale well when application demand increases. It can handle more data by adding more nodes to the cluster. Based on tests [9] conducted

on RabbitMQ with set of 4K(4096) and 16K (16384) messages, the performance of RabbitMQ on multi node cluster decreases in comparison with single node cluster to reach a threshold and then becomes stable. This decrease in performance can be primarily accounted due to replication between nodes in the cluster. These tests were conducted on combination of single publisher and single subscriber, multiple publishers and single subscriber, multiple publishers and multiple subscribers. Due to time constraints there is no concrete conclusion to these test and numbers. Further testing should be conducted on a bigger cluster i.e. more nodes and for longer duration to produce more consistent results.

ACKNOWLEDGEMENTS

Special thanks to Professor Gregor von Laszewski, Dimitar Nikolov and all associate instructors for all help and guidance related to latex and bibtex, scripts for building the project, quick and timely resolution to any technical issues faced. The paper is written during the course I524: Big Data and Open Source Software Projects, Spring 2017 at Indiana University Bloomington.

REFERENCES

- [1] Pivotal, "RabbitMQ, clients and developer tools," Web Page, accessed: 2017-02-26. [Online]. Available: <https://www.rabbitmq.com/devtools.html>
- [2] Pivotal, "RabbitMQ, components," Web Page, accessed: 2017-02-02. [Online]. Available: <https://www.rabbitmq.com/>
- [3] S. Vinoski, "Advanced message queuing protocol," *IEEE Internet Computing*, vol. 10, no. 6, 2006. [Online]. Available: <http://ieeexplore.ieee.org/stamp/stamp.jsp?arnumber=4012603>
- [4] Wikipedia, "Producer-consumer problem," Web Page, accessed: 03-18-2017. [Online]. Available: https://en.wikipedia.org/wiki/Producer%E2%80%93consumer_problem
- [5] A. Videla and J. J. Williams, *RabbitMQ in action*. Manning, 2012.
- [6] Lovisa Johansson, "Rabbitmq for beginners - what is rabbitmq?" Web Page, May 2015, accessed: 02-15-2017. [Online]. Available: <https://www.cloudamqp.com/blog/2015-05-18-part1-rabbitmq-for-beginners-what-is-rabbitmq.html>
- [7] Pivotal, "RabbitMQ license," Web Page, accessed: 2017-02-02. [Online]. Available: <https://www.rabbitmq.com/mpl.html>
- [8] Wikipedia, "Mozilla public license," Web Page, accessed: 2017-03-18. [Online]. Available: https://en.wikipedia.org/wiki/Mozilla_Public_License
- [9] B. Jones, S. Luxenberg, D. McGrath, P. Trampert, and J. Weldon, "Rabbitmq performance and scalability analysis," *project on CS*, vol. 4284, 2011. [Online]. Available: <https://people.cs.vt.edu/butta/cs4284/spring2011/butta/RabbitMQPaper.pdf>