

Oracle PGX

PIYUSH SHINDE^{1,*}

¹ School of Informatics and Computing, Bloomington, IN 47408, U.S.A.

* Corresponding authors: piyushssinde1992@gmail.com

Paper-1, May 3, 2017

Oracle PGX is a graph analysis toolkit. Graph analysis provides key insights, which play a vital role in business optimizations. Oracle PGX provides various graph analysis features like loading graphs into memory, running built-in and user-defined graph analysis algorithms, running graph pattern matching queries and mutating graphs [1]. In this paper we discuss Oracle PGX's various features, use cases, advantages and limitations.

© 2017 <https://creativecommons.org/licenses/>. The authors verify that the text is not plagiarized.

Keywords: Graph Analysis, Algorithms, Shell

<https://github.com/cloudmesh/sp17-i524/blob/master/paper1/S17-IR-2035/report.pdf>

1. INTRODUCTION

A graph is a fundamental data structure that defines links between different data nodes. Graphs represent data sets in a wide range of applications. In a social graph, for example, nodes correspond to people while friendship relationships between them are represented as edges [2].

Graph analysis involves extracting information from a dataset, which is represented as a graph. Large amounts of data demands large amounts of computational power to analyse. Challenges faced during graph analysis include storage capacity for large data sets, performance variation with varying data set sizes, and efficient implementation of graph analyses algorithms.

Oracle PGX (Parallel Graph AnalytiX) is a fast, parallel, in-memory graph analytic framework. Its main features include loading graphs into memory, running built-in and user-defined graph analysis algorithms, running graph pattern matching queries and mutating graphs [1]. We will glance through these features in this paper.

Oracle PGX 2.3.1 is the latest version available as of 26th February 2017 [3].

2. INSTALLATION

To install Oracle PGX the platform requirements include a Linux or Mac OS X platform [3].

Oracle PGX can be installed in the following steps. Download Java SE 8 Development Kit (JDK8). Make sure the JAVA_HOME environment variable points to the JDK8 home directory, e.g. export JAVA_HOME=/usr/lib/jvm/java-8-oracle. Unzip the downloaded zip file e.g. unzip pgx-2.3.1-server.zip -d /opt/pgx. Change directory (cd) into the installation directory. e.g. cd \$PGX_HOME. Start the PGX shell using the command ./bin/pgx.

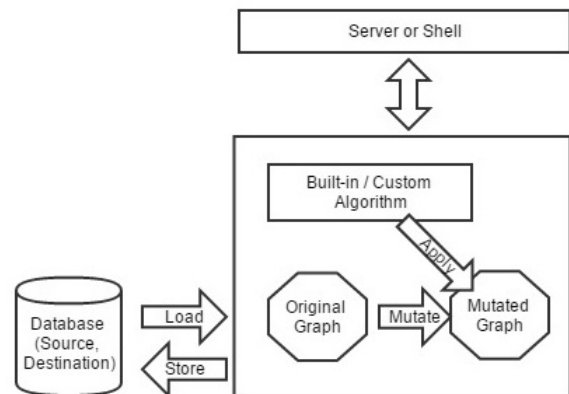


Fig. 1. PGX Overview [1]

3. USAGE MODES

PGX provides four different usage modes, depending on different user requirements [4].

3.1. Local Shell Mode

This mode enables the user to interact with a PGX shell on a local machine through a command-line shell interface. User commands are dynamically interpreted and executed by the PGX shell to invoke the PGX function. It can be used to run some quick graph analyses locally on small graph data. It can likewise be used to try different graph algorithms to explore the different interpretations of a data set.




Fig. 2. PGX Local Shell Mode [4]

3.2. Remote Shell Mode

In the remote shell mode, the PGX execution engine is deployed as a RESTful application on a powerful PGX-Server which is remotely connected to the PGX-Shell from a client machine. This mode is useful for performing graph analysis on large data sets. This mode enables to connect multiple clients to the PGX-Server.




Fig. 3. PGX Remote Shell Mode [4]

3.3. Local Java Application Mode

This modes allows embedding PGX into a Java application, and using graph analysis as a part of the application's functionality.




Fig. 4. PGX Locally Embedded Mode [4]

3.4. Remote Java Application Mode

This mode empowers the user to use the Java application to control the PGX-Server remotely, from a client machine. Similar to the remote shell model, this mode enables connections to the PGX-Server via multiple clients.

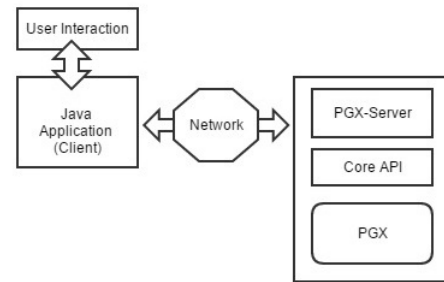


Fig. 5. PGX Remotely Embedded Mode.

4. GRAPH MUTATIONS

PGX helps users create mutated versions of the original graph, to perform algorithms for graph analyses. Since PGX does not allow mutation of the original graph, it always creates a new mutated version of the original graph.

A graph can be simplified by removing self-edges, duplicated edges, and trivial vertices.




Fig. 6. Simplifying a graph. Source: Paper [5]

Graphs can also be simplified by undirecting it's edges. Sub-graphs and bipartite subgraphs can also be created in PGX.

5. RUNNING GRAPH ALGORITHMS ON PGX

PGX allows you to run graph algorithms in two ways.

Built-in graphs: PGX provides a set of built-in graph algorithms that are ready to use. They include computing various centrality measures, finding shortest paths, finding/evaluating clusters and components, and predicting future edges, etc. The detailed and complete list of the built-in algorithms, with their time complexity, space requirement and implementation code is available on the official website of Oracle Labs PGX.

Custom algorithms: User-defined algorithms can be programmed in PGX (excluding the ones from the built-in package), using the Green-Marl language (a Domain Specific Language) [6]. The Domain Specific Language approach provides benefits like productivity, performance and portability.

6. USE CASES

PGX can be used to solve real world graph analyses problems. Here are two such applications [7].

6.1. Detect Potential Fraud in a Healthcare System

In this example, anomalies in medical transactions were detected using PGX.

A public data set of medical transactions from CMS (United States Center for Medicare and Medicaid Services) for the year 2012 was used. It had transactions between over 880,000 medical providers and CMS, adding up to more than \$77B for 2012 [8].

To analyse the data, the data set was converted to a bipartite graph, separating health care providers from health services. An edge in the bipartite graph linked a specific health service to a health care provider. A two-hop (undirected) path existed between two health care providers, for each common health service they both provided. The more common services between two health care providers, the closer they would be to each other in the graph. The data set also specified the ‘specialities’ of medical providers. Vertices of medical providers of the same speciality would be closer to each other, since they would provide similar services. In case, a medical provider was exceptionally close to other providers with different speciality, it would be considered an anomaly since the medical provider would be providing services unfitting to its speciality.

Personalized PageRanks (PPR) were used to compute closeness between vertices. For each speciality a PPR score was computed. These scores were then cross-checked to see if they belonged to the current speciality. If not, the provider was considered an anomaly. This graph analysis was implemented using the built-in Personalized PageRank algorithm.

The results were further investigated to identify potential frauds in the Healthcare System.

6.2. Using PGX as a Recommendation Engine

This example helps us discover the information implied by relationships through graph analytics.

A publicly available data set ‘MovieLens’ was used. Recommendations were generated using matrix factorization algorithm. Matrix factorization is a simple process, in which you factorize a matrix into two matrices, such that the newly generated matrices when multiplied with each other would generate back the original matrix [9].

We had a matrix of movies and users, where the values were the ratings given by the user to a movie. The aim was to recommend a bunch of movies for each user.

An array of features was synthesized, considering user information, the movies they rated and their ratings. The approach was to initialize each matrix with some values. Each user had an array of floating-point values that represented how strongly they were associated with that feature. To generate recommendations, the movie’s feature array was multiplied to the user’s feature array. The corresponding values were then sorted in a descending fashion and the top few items were the recommendations. This process was very fast using PGX. The results were then used for many recommendations, repeating the process periodically after new users or movies were added to the system.

7. BENEFITS AND LIMITATIONS

PGX executes graph algorithms in a fast, parallel, in-memory fashion. It provides built-in algorithms and supports user-defined algorithms. It offers an interleaved usage of graph analysis algorithms and graph pattern-matching. In addition, it provides an interactive shell application to easily execute commands from the shell command line.

PGX loads file-systems, exclusively in graph formats. It requires licensing, restricting its usage. Both of the above mentioned limitations can be removed by contacting Oracle Labs PGX [10].

8. USEFUL RESOURCES

The complete documentation of Oracle PGX, including instructions to download and install and several tutorials, in a systematic method are available on the main website of Oracle Labs PGX [11].

The paper titled “Graph Analysis – Do We Have to Reinvent the Wheel?”, discusses about the alternative database technologies that provide optimized performance with regards to the dedicated graph databases. [12]

The paper on Green-Marl DSL approach titled “Green-Marl: A DSL for Easy and Efficient Graph Analysis”, provides insight on the methods of using the Green-Marl DSL technique for optimizing graph analyses using user-defined algorithms. [2]

9. CONCLUSION

Numerous information is revealed from graphs. The analysis of graphs can unveil significant insights. Oracle PGX (Parallel Graph AnalytiX) is a toolkit for graph analysis. It allows users to load up their graph data and run analytic algorithms on them. It provides different usage modes which can be used as per client requirements. It allows users to implement various built-in and user-defined graphs and mutate graphs. Some classes of the built-in algorithms are community detection, path finding, page ranking, recommendation, pattern matching and influencer identification. Oracle PGX is easy to use. It can be used for precise and efficient graph analysis, and is scalable to implement graph analysis algorithms for large data sets.

10. ACKNOWLEDGMENTS

This project was a part of the Big Data and Software and Projects (INFO-I524) course. I would like to thank Professor Gregor von Laszewski and the associate instructors for their help and support during the course.

REFERENCES

- [1] Oracle Labs, “PGX Overview,” Web Page, accessed: 2017-02-26. [Online]. Available: https://docs.oracle.com/cd/E56133_01/1.2.0/reference/overview/index.html
- [2] S. Hong, H. Chafi, E. Sedlar, and K. Olukotun, “Green-marl: a dsl for easy and efficient graph analysis,” in *ACM SIGARCH Computer Architecture News*, vol. 40, no. 1. ACM, 2012, pp. 349–362.
- [3] Oracle Labs, “Oracle Labs PGX Downloads,” Web Page, accessed: 2017-02-24. [Online]. Available: <http://www.oracle.com/technetwork/oracle-labs/parallel-graph-analytics/downloads/index.html>
- [4] Oracle Labs, “PGX Usage Modes,” Web Page, accessed: 2017-02-24. [Online]. Available: https://docs.oracle.com/cd/E56133_01/2.3.1/reference/overview/usage.html
- [5] Oracle Labs, “Mutating Graphs,” Web Page, accessed: 2017-02-24. [Online]. Available: https://docs.oracle.com/cd/E56133_01/2.3.1/reference/overview/mutate.html
- [6] Oracle Labs, “Running Graph Algorithms on PGX,” Web Page, accessed: 2017-02-24. [Online]. Available: https://docs.oracle.com/cd/E56133_01/2.3.1/reference/overview/run.html
- [7] Oracle Labs, “Use Cases,” Web Page, accessed: 2017-02-26. [Online]. Available: https://docs.oracle.com/cd/E56133_01/1.2.0/use-cases/index.html
- [8] Oracle Labs, “Detect Anomalies in a Healthcare System,” Web Page, accessed: 2017-02-26. [Online]. Available: https://docs.oracle.com/cd/E56133_01/1.2.0/use-cases/healthcare-fraud/index.html
- [9] Oracle Labs, “Using PGX as a Recommendation Engine,” Web Page, accessed: 2017-02-26. [Online]. Available: https://docs.oracle.com/cd/E56133_01/1.2.0/use-cases/recommendation/index.html

- [10] Oracle Labs, "About This OTN Release," Web Page, accessed: 2017-02-24. [Online]. Available: https://docs.oracle.com/cd/E56133_01/2.3.1/about/index.html
- [11] Oracle Labs, "Oracle Labs PGX," Web Page, accessed: 2017-02-24. [Online]. Available: https://docs.oracle.com/cd/E56133_01/1.2.0/index.html
- [12] A. Welc, R. Raman, Z. Wu, S. Hong, H. Chafi, and J. Banerjee, "Graph analysis: do we have to reinvent the wheel?" in *First International Workshop on Graph Data Management Experiences and Systems*. ACM, 2013, p. 7.