

Berkeley DB

SABER SHEYBANI^{1,*}

¹School of Informatics and Computing, Bloomington, IN 47408, U.S.A.

* Corresponding authors: sheybani@indiana.edu

Paper 2, April 30, 2017

Berkeley DB is a family of open source, NoSQL key-value database libraries. It provides a simple function-call API for data access and management over a number of programming languages, including C, C++, Java, Perl, Tcl, Python, and PHP. Berkeley DB is embedded because it links directly into the application and runs in the same address space as the application. As a result, no inter-process communication, either over the network or between processes on the same machine, is required for database operations. It is also extremely portable and scalable, it can manage databases up to 256 terabytes in size. For data management, Berkeley DB offers advanced services, such as concurrency for many users, ACID transactions, and recovery. Berkeley DB is used in a wide variety of products and a large number of projects, including gateways from Cisco, Web applications at Amazon.com and open-source projects such as Apache and Linux.

© 2017 <https://creativecommons.org/licenses/>. The authors verify that the text is not plagiarized.

Keywords: NoSQL, embedded database, Oracle, open source database management system

<https://github.com/cloudmesh/sp17-i524/tree/master/paper2/S17-ER-1001/report.pdf>

1. INTRODUCTION

Data management has always been a fundamental issue in programming. Since 1960s, countless database management systems have been developed to fulfil different sorts of demands. The question for every user is choosing the system that best fits the requirements of its application.

Database management systems can be categorized based on data models, into a number of groups: Hierarchical Databases, Network Databases, Relational Databases, Object-based Databases, and Semistructured Databases.

Hierarchical databases use the oldest type of data models, which is a tree-like structure. The records are connected to each other with a hard-coded link.

In a **Network** model, records are also connected with links, but there is no hierarchy. Instead, the structure is graph-like and all of the nodes can connect to each other.

In **Relational** databases, there are no physical links, but the data is structured in tables (relations). Each row represents a record and each column represents an attribute. The tables are connected with common attributes, which makes querying much easier than the two former models. For this reason, database management systems using relational models are the most widely used ones.

Object-based data models extend concepts of object-oriented programming into database systems, in order to provide persistent storage of objects and other capabilities of databases for object-oriented programming.

Semistructured databases which include NoSQL databases are the type of database model that enable storage of heterogeneous data, by allowing records with different attributes. This however, is achieved by sacrificing the knowledge of data type by the database system. The data in this case must be *self-describing*, meaning that the description (schema) of the data must be in itself. XML (Extensible Markup Language) schema language is a widely used language for providing schema for these database systems[1].

Berkeley DB fits into the last category, as a NoSQL database system. The records are stored as key-value pairs and a few logical operations can be executed on them, namely: insertion, deletion, finding a record by its key, and updating an already found record. "Berkeley DB never operates on the value part of a record. Values are simply payload, to be stored with keys and reliably delivered back to the application on demand." There is no notion of schema and no support for SQL queries. "The application must understand the keys and values that it uses. On the other hand, there is literally no limit to the data types that can be stored in a Berkeley DB database. The application never needs to convert its own program data into the data types that Berkeley DB supports. Berkeley DB is able to operate on any data type the application uses, no matter how complex" [2].

2. ARCHITECTURE

Berkeley DB's architecture can be explained by five major sub-systems: **Access Methods:** Providing general-purpose support

for creating and accessing database files. **Memory Pool:** The general-purpose shared memory buffer pool. Multiple **Transaction:** Implementing the transaction model, realizing ACID properties. processes and threads within processes share access to databases using this subsystem. **Locking:** The general-purpose lock manager for processes. **Logging:** The write-ahead logging that supports the Berkeley DB transaction model.

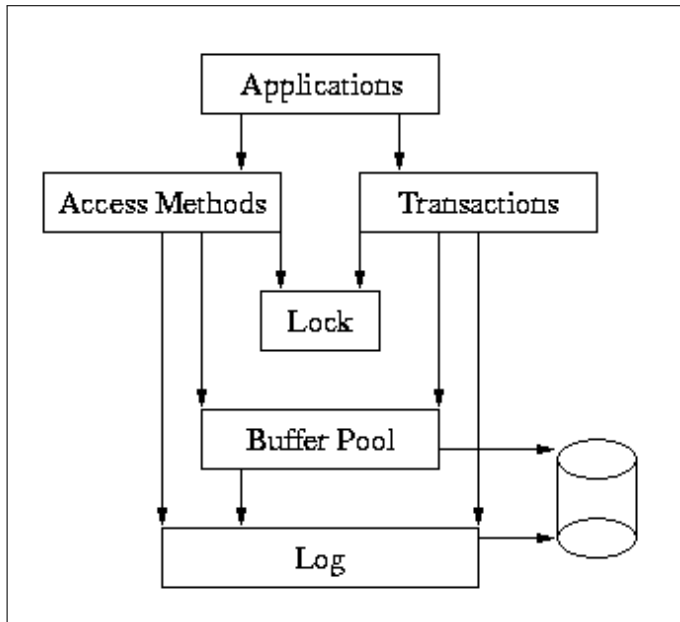


Fig. 1. Berkeley DB Subsystems [3]

Figure 1 displays a diagram of the Berkeley DB library architecture. The arrows are calls that invoke the destination. Each subsystem can also be used independent from the other ones, but this usage is not common.

3. SERVICES AND OTHER FEATURES

The two fundamental services that every database management system provides are data access, and data management services. Data access services include the low-level operations on the records, which were already mentioned in the introduction for the case of Berkeley DB. In terms of storage structure, Berkeley DB supports **hash tables, Btrees, simple record-number-based storage, and persistent queues** [4].

Data management services are the higher-level services (and features) such as concurrency that ensure specific qualities for operation of the system. These services include allowing simultaneous access to the records by multiple users (concurrency), changing multiple records at the same time (transaction), and complete recovery of the data from crashes (recovery) [4].

For **concurrency**, Berkeley DB is able to handle low-level services such as locking and shared buffer management transparently, while multiple processes and threads use the data. For recovery, every application can ask Berkeley DB for recovery, at startup time.

An **ACID transaction** ensures the following specifications at the end of its operation [5]: Atomicity (Either all or none of the records change), Consistency (The system goes from one valid state to another), Isolation (concurrent execution of multiple transaction yields the same result as the sequential execution of them), Durability (The result remains steady, even in case of

crash of the system). Berkeley DB "libraries provide strict ACID transaction semantics, by default. However, applications are allowed to relax the isolation guarantees the database system makes" [4].

Berkeley DB runs in the same address space as the application. As a result, there is no need for communication between processes and threads. On the other hand, as an embedded database management system, it does not provide a standalone server. However, server applications can be built over Berkeley DB and many examples of Lightweight Directory Access Protocol (LDAP) servers have been built using it [2].

The database library for Berkeley DB consumes less than 300 kilobytes of text space on common architectures. That makes it a feasible solution for embedded systems with small capacities. Nonetheless, it can manage up to 256 terabytes databases.

3.1. Supported Operating Systems and Languages

Berkeley DB supports nearly all modern operating systems. They include Windows, Linux, Mac OS X, Android, iPhone, Solaris, BSD, HP-UX, AIX, and RTOS such as VxWorks, and QNX. The supported programming languages include "C, C++, Java, C#, Perl, Python, PHP, Tcl, Ruby and many others" [6].

3.2. Required Infrastructure

As infrastructure, Berkeley DB requires "underlying IEEE/ANSI Std 1003.1 (POSIX) system calls and can be ported easily to new architectures by adding stub routines to connect the native system interfaces to the Berkeley DB POSIX-style system calls" [7].

4. PRODUCTS AND LICENSING

The products include three implementations on C, C++, and Java (Oracle Berkeley DB, Oracle Berkeley XML, and Oracle Berkeley JE, respectively) [8].

Berkeley DB is an open source library and is free for use and redistribution in other open source products. the distribution includes complete source code for all three implementations, their supporting utilities, as well as complete documentation in HTML format [7].

For redistribution in commercial products, Sleepycat Software licenses four products, with prices ranging from US\$900 to 13,800 per processor [9] as of March 2017. The products, in the order of ascending price and capabilities are: Berkeley DB Data Store, Berkeley DB Concurrent Data Store, Berkeley DB Transactional Data Store, Berkeley DB High Availability. The Sleepycat software also includes prebuilt libraries and binaries as part of support services, which is not provided in the free distribution. There is no additional license payment for embedded usage within the Oracle Retail Predictive Application Server (RPAS).

5. USE CASES

A notable number of open source and commercial products in different areas of technology, use Berkeley DB. Open source use cases include Linux, UNIX, BSD, Apache, Solaris, MySQL, Sendmail, OpenLDAP, and MemcacheDB.

Proprietary applications "include directory servers from Sun and Hitachi; messaging servers from Openwave and LogicaCMG; switches, routers and gateways from Cisco, Motorola, Lucent, and Alcatel; storage products from EMC and HP; security products from RSA Security and Symantec; and Web applications at Amazon.com, LinkedIn and AOL" [6].

6. ADVANTAGES AND LIMITATIONS

Berkeley DB has two advantages over relational and object-oriented database systems, when it comes to embedded applications. One is running in the same address space as the application and thus, not requiring any inter-process communication which can have a high cost in embedded applications. And the other is simplicity of interface for operations which does not require query language parsing. These two features along with its small size, give Berkeley DB system a privilege of being lightweight enough for many applications where there are tight constraints on resources.

However, with simplicity comes the lack of SQL features. If the user of the application needs to perform complicated searches (potentially using SQL queries) the programmer would need to write the code for those cases. In general, Berkeley DB is aimed at providing fast, reliable, transaction-protected record storage, at a minimalist way [10].

7. EDUCATIONAL MATERIAL

As was mentioned in the Products section, the free distribution comes with complete documentation in HTML format. The documentation has two parts: a reference manual in UNIX-style for programmers, and a reference guide which can serve as a tutorial [11]. In addition to that, *Berkeley DB Tutorial and Reference Guide, Version 4.1.24* [7] and *The Berkeley DB Book* [12] are useful resources for learning more about Berkeley DB and getting started with it.

8. CONCLUSION

Berkeley DB is a minimal, lightweight database management system, focused on providing performance, especially in embedded systems. It offers a small, simple set of data access services, and a rich powerful set of data management services. It is freely available for use by non-commercial distributions and has been successfully used in many projects.

REFERENCES

- [1] I. Limited, *Introduction to Database Systems*. Pearson, 2010. [Online]. Available: <https://books.google.com/books?id=-YY-BAAAQBAJ>
- [2] "What berkeley db is not," Web Page, accessed: 2017-4-6. [Online]. Available: <https://web.stanford.edu/class/cs276a/projects/docs/berkeleydb/ref/intro/dbisnot.html>
- [3] "The big picture," Web Page, accessed: 2017-4-6. [Online]. Available: <https://web.stanford.edu/class/cs276a/projects/docs/berkeleydb/ref/arch/bigpic.html>
- [4] "What is berkeley db?" Web Page, accessed: 2017-4-6. [Online]. Available: <https://web.stanford.edu/class/cs276a/projects/docs/berkeleydb/ref/intro/dbis.html>
- [5] T. Haerder and A. Reuter, "Principles of transaction-oriented database recovery," *ACM Computing Surveys (CSUR)*, vol. 15, no. 4, pp. 287–317, 1983.
- [6] "Oracle berkeley database products," Web Page, accessed: 2017-4-6. [Online]. Available: <http://www.oracle.com/technetwork/products/berkeleydb/learnmore/berkeley-db-family-datasheet-132751.pdf>
- [7] "Berkeley db tutorial and reference guide, version 4.1.24," Web Page, accessed: 2017-4-6. [Online]. Available: <https://web.stanford.edu/class/cs276a/projects/docs/berkeleydb/reftoc.html>
- [8] "Oracle berkeley db 12c," Web Page, accessed: 2017-4-6. [Online]. Available: <http://www.oracle.com/technetwork/database/database-technologies/berkeleydb/%20overview/index.html>
- [9] "Oracle technology global price list," Web Page, accessed: 2017-4-6. [Online]. Available: <http://www.oracle.com/us/corporate/pricing/technology-price-list-070617.pdf>

- [10] "Do you need berkeley db?" Web Page, accessed: 2017-4-6. [Online]. Available: <https://web.stanford.edu/class/cs276a/projects/docs/berkeleydb/ref/intro/need.html>
- [11] M. A. Olson, K. Bostic, and M. I. Seltzer, "Berkeley db." in *USENIX Annual Technical Conference, FREENIX Track*, 1999, pp. 183–191.
- [12] H. Yadava, *The Berkeley DB Book*, ser. Books for Professionals by Professionals. Apress, 2007. [Online]. Available: <https://books.google.com/books?id=2wEkW7pQ0KwC>

AUTHOR BIOGRAPHY

Saber Sheybani received his B.S. (Electrical Engineering - Minor in Control Engineering) from University of Tehran. He is currently a PhD student of Intelligent Systems Engineering - Neuroengineering at Indiana University Bloomington.