

Apache Derby

RIBKA RUFAEL^{1,*}, +

¹School of Informatics and Computing, Bloomington, IN 47408, U.S.A.

* Corresponding authors: rrufael@umail.iu.edu

+HID: S17-IO-3016

paper2, April 30, 2017

Apache Derby, part of Apache DB subproject, is Java based relational database management system. Apache Derby database provides storage, access and secure management of data for Java based applications. Apache Derby is an open source software and it is licensed under Apache version 2.0. Apache Derby is written in Java and it runs on any certified JVM(Java Virtual Machine). JDBC driver in embedded or network server frameworks allows applications to access Apache Derby Database.

© 2017 <https://creativecommons.org/licenses/>. The authors verify that the text is not plagiarized.

Keywords: Apache Derby, relational database management system, JDBC

<https://github.com/cloudmesh/sp17-i524/raw/master/paper2/S17-IO-3016/report.pdf>

1. INTRODUCTION

Apache Derby is an open source relational database management system(RDBMS). Applications interact with Apache Derby database through JDBC(Java Database Connectivity) driver. Apache Derby is built in Java programming language which makes it platform independent. Apache Derby can be embedded into Java based application or it can be setup to run in a client/server environment. Apache Derby complies with JDBC and ANSI SQL standards. Apache Derby database allows applications to create, update, read, delete and manage data [1, 2].

Apache Derby is part of the Apache DB subproject of Apache Software Foundation. It was in August 2004 that IBM submitted the Derby code to Apache Software Foundation. Then Apache Derby became part of the Apache DB project in July 2005 [2].

2. ARCHITECTURE

There are two views that describe Apache Derby's architecture. These views are Module View and Layer/Box view [3].

2.1. Module View

Modules and Monitor are components in Apache Derby database system. A collection of distinct functionality is a module. Examples of modules in Apache Derby are lock management, error logging and JDBC driver. Lock management is responsible for controlling concurrent transactions on data objects. Once Apache Derby database is up and running any informational messages or error messages are logged. This message logging is handled by an error logging module. Applications use JDBC driver to interact with Apache Derby database. A number of classes are used for the implementation of each module [3].

The monitor is responsible for managing Apache Derby database. Whenever a request to modules come, the monitor is responsible for selecting appropriate module implementation depending on what the module request was and from which environment the request came from [3].

2.2. Layer/Box View

JDBC, SQL, Store and Services are the four layers in Apache Derby. The Java Database Connectivity abbreviated JDBC is an API(Application Programming Interface) which allows applications to connect to Apache Derby database. JDBC interfaces that Apache Derby implements allow applications to connect to Apache Derby. Some of the interfaces are Driver, DataSource, ConnectionPoolDataSource, XADataSource and PreparedStatement. Apache Derby JDBC has implementation of java.sql and javax.sql classes [3].

The other layer which sits below JDBC is SQL layer. Compilation and execution are the two logical parts of SQL layer. SQL statement that is invoked by an application to Apache Derby Database passes through five step compilation process. First statement is parsed with a parser created by JavaCC(Java Compiler Compiler) and tree of query nodes is created. Second step is binding to resolve objects. Third step is optimizing to identify the access path. In fourth step Java class is created for the statement this is then cached to be used by other connections. Finally, on the fifth step, class is loaded and instance of generated statement is created. During execution, execute methods on the class instance created during compilation are called and Derby result set is returned. JDBC layer is responsible in converting the Derby result set into JDBC result set for the applications [3].

Access and raw are the two parts of the store layer. Raw data storage for data in files and pages, transaction logging

and management is handled by the raw store. The access store interfaces with SQL layer and takes care of scanning of tables and indexes, indexing, sorting, locking and etc [3].

Lock management, error logging and cache management are part of the service layer. Clock based algorithm is used by cache management. It is mainly used for caching buffer, caching compiled java classes for SQL statement implementation plans and caching table descriptors [3].

2.3. Shell Access

Shell access to Apache Derby database is achieved by ij. The ij tool is one of java utility tools that allows performing sql scripts on Derby database in embedded or network server frameworks. The ij tool can be used for creating database, connecting to database, run and execute sql scripts. Commands in ij are case sensitive and semicolon is used to mark end of command. Other utility tools that come with Apache Derby are sysinfo, dblook and SignatureChecker. The sysinfo utility tool is used to get version and other information about Apache Derby and the Java environment. The dblook utility is used to generate Data Definition Language (DDL) for Derby database. Checks, functions, indexes, jar files, primary keys, foreign keys and schemas are the objects generated by dblook. SignatureChecker is used to check whether functions and procedures used in Derby database comply with the rules and standards [4].

2.4. API

Applications can interact with Apache Derby database through two JDBC drivers `org.apache.derby.jdbc.EmbeddedDriver` and `org.apache.derby.jdbc.ClientDriver` for embedded and network server frameworks respectively [4].

3. INSTALLATION

Java 2 Standard Edition (J2SE) 6 or higher is needed for installing Apache Derby and for running Derby the Java Runtime Environment (JRE) is needed. Apache Derby can be downloaded from the Apache DB download page [5]. Apache Derby can be installed on Windows, MAC, UNIX and Linux operating systems. After downloading the zipped installation file, the file should be extracted into directory and then `DERBY_INSTALL` variable should be set in the same directory as where Derby was installed. Apache Derby provides two frameworks Embedded Derby and Derby Network Server [6].

3.1. Embedded Derby

In the embedded Derby framework, Apache Derby engine and the application which accesses it run on the same JVM (Java Virtual Machine). Embedded Derby JDBC driver is used by the application to interact with Apache Derby database. To setup Embedded Derby mode, `derby.jar` and `derbytools.jar` must be included in the CLASSPATH after installation. The Derby engine and Embedded Derby JDBC driver are included in `derby.jar`. `derbytools.jar` includes ij tool (utility tool which can be used as scripting tool to interact with Derby database). In Embedded Derby framework, multiple users that are running in the same JVM can access the same database. Figure 1 shows the Embedded Derby framework [6].

3.2. Derby Network Server

Derby Network Server framework allows multiple application running in the same JVM or different JVMs to access Derby




Fig. 1. Embedded Derby Framework [6].

database over the network in a typical client/server architecture. In this framework application accesses Derby database through Derby Network Client JDBC driver. To setup Derby network server `derbynet.jar` and `derbytools.jar` must be included in CLASSPATH on the server side. The program for Network Server and reference to the Derby engine are included `derbynet.jar` file. On the client side, `derbyclient.jar` and `derbytools.jar` must be included in CLASSPATH. Derby Network Client JDBC driver is included in `derbyclient.jar` file. Derby Network Server listens and accepts requests on port 1527 by default but this can be changed to a different port if needed. Figure 2 shows the Derby Network Server framework [6].

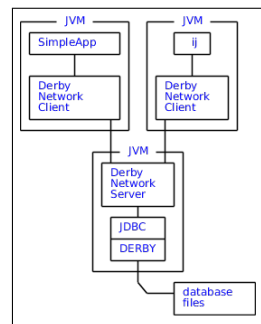


Fig. 2. Derby Network Server Framework [6].

Another variation for setting up Derby Network Server is embedded server. In embedded server, application will have both Embedded Derby JDBC driver and Network Server. The application uses Embedded Derby JDBC driver which runs on the same JVM and extends access to other applications running on different JVMs through the Network Server [6].

4. SECURITY

Apache Derby provides a number of security options. Some of these features are authentication, authorization and disk encryption. Before users are granted access to the Derby database, Derby can be setup to perform user authentication. There may be a need for certain user groups to have read only access to Derby database and some other user groups to have both read and write access to Derby database which can be achieved by Derby's user authorization feature. Data which is saved on disk can be encrypted with Derby's disk encryption feature [7].

5. USE CASES

Apache Hive, data warehouse querying and analysis software, uses Apache Derby database by default for metadata store. Apache Derby can be configured in the embedded or network server mode for Hive meta data storage [8, 9].

My Money, management and analysis software for personal and business finances which is built by MTH Software, Inc, uses Apache Derby as a relational database management system [10, 11].

6. LICENSING

Apache Derby is an open source technology hence it is available for free. Apache Derby is licensed under Apache License, Version 2.0 [1].

7. EDUCATIONAL MATERIAL

Apache Derby tutorial page is one of the resources which can be used by users who are new to Apache Derby. This tutorial contains step by step information about Apache Derby installation, embedded framework configuration and network server framework configuration [6]. Derby Engine Architecture Overview page provides information about Apache Derby architecture [3]. Apache Derby documentation page has list of documentations that can be used as reference for new users, developers and administrators [12].

8. CONCLUSION

Apache Derby database offers storage, access and secure management of data for Java based applications. Apache Derby provides complete relational database management package that complies with JDBC and ANSI SQL standards. Apache Derby's small size makes it suitable for embedding it into applications which run on smaller devices with less physical memory. As the use case of Hive using Apache Derby for metadata storage indicates, Apache Derby can be incorporated into software stack for big data projects for metadata storage.

ACKNOWLEDGEMENTS

The author would like to thank Professor Gregor von Laszewski and associate instructors for their help and guidance.

REFERENCES

- [1] Apache Software Foundation, "Apache Derby," Web Page, Oct. 2016, accessed: 2017-03-14. [Online]. Available: <https://db.apache.org/derby/>
- [2] The Apache Software Foundation, "Apache Derby Project Charter," Web page, Sep. 2016, accessed: 2017-02-25. [Online]. Available: https://db.apache.org/derby/derby_charter.html
- [3] Apache Software Foundation, "Derby Engine Architecture Overview," Web Page, Sep. 2016, accessed: 2017-03-14. [Online]. Available: http://db.apache.org/derby/papers/derby_arch.html#Module+View
- [4] Apache Software Foundation, "Derby Tools and Utilities Guide," Web Page, Oct. 2016, accessed: 2017-03-24. [Online]. Available: <https://db.apache.org/derby/docs/10.13/tools/derbytools.pdf>
- [5] Apache Software Foundation, "Apache Derby: Downloads," Web Page, Sep. 2016, accessed: 2017-03-24. [Online]. Available: http://db.apache.org/derby/derby_downloads.html
- [6] Apache Software Foundation, "Apache Derby Tutorial," Web Page, Sep. 2016, accessed: 2017-03-14. [Online]. Available: <https://db.apache.org/derby/papers/DerbyTut/index.html>
- [7] Apache Software Foundation, "Derby and Security," Web Page, Mar. 2017, accessed: 2017-03-20. [Online]. Available: <http://db.apache.org/derby/docs/10.8/devguide/cdevcsecuree.html>
- [8] Apache Software Foundation, "Apache Hive," Web Page, Mar. 2017, accessed: 2017-03-23. [Online]. Available: <https://cwiki.apache.org/confluence/display/Hive/Home>
- [9] Apache Software Foundation, "GettingStarted-Metadata Store," Web Page, Mar. 2017, accessed: 2017-03-23. [Online]. Available: <https://cwiki.apache.org/confluence/display/Hive/GettingStarted#GettingStarted-MetadataStore>
- [10] MTH Software, Inc, "MTH built products," Web Page, Mar. 2017, accessed: 2017-03-23. [Online]. Available: <http://www.mthbuilt.com/products.php>
- [11] MTH Software, Inc, "How to use SQL in My Money," Web Page, Mar. 2017, accessed: 2017-03-23. [Online]. Available: http://wiki.mthbuilt.com/How_to_use_SQL_in_My_Money
- [12] Apache Software Foundation, "Apache Derby: Documentation," Web Page, Oct. 2016, accessed: 2017-03-24. [Online]. Available: <https://db.apache.org/derby/manuals/index.html>