

Remotely Deploying, Visualizing and Controlling a Robot Swarm with ROS and Cloudmesh

MATTHEW LAWSON¹ AND GREGOR VON LASZEWSKI^{1,*}

¹ School of Informatics and Computing, Bloomington, IN 47408, U.S.A.

* Corresponding authors: laszewski@gmail.com

Project: S17-IO-3010, September 21, 2017

Working with robot swarms has emerged as an important area of robotics research. Since the cost of deploying physical robot swarms could prove time-consuming and cost-prohibitive to research departments, the field needs a convenient, cost-efficient method of educating the next generation of experts. Hence, we explore the feasibility of simulating such swarms while harnessing remotely-located, distributed computing environments to conduct such simulations. In particular, our interest lies in using and re-using cloud computing resources to simulate large-scale robot swarms. Therefore, we produced the proof-of-concept deployment model described in the following pages, which creates a two-robot virtual swarm on multiple clouds to which we have access. Importantly, deployment occurs with minimal user intervention, opening up the research field to STEM students and undergraduate freshmen, as well as PhD students. Users can observe the robots' behavior through a streamed visual simulation. We demonstrate that users can start up the behavior of a robot through a single command, allowing us to benchmark the activity streams of the swarm in the cloud. The project uses open source software throughout, including the Open Source Robotics Foundation's (OSRF) Robot Operating System to define, create and control the virtual robots. The OSRF's *Gazebo* program provides the visualization of the simulation. *Cloudmesh* simplifies interaction with the various clouds we use.

Keywords: Cloud, I524, ROS, Gazebo, Robot, Swarm

Report: <https://github.com/cloudmesh/sp17-i524/tree/master/project/S17-IO-3010/report/report.pdf>

Code: <https://github.com/cloudmesh/cloudmesh.ros>

1. INTRODUCTION

With the evolution of robotics and its increased influence in society we anticipate the need to deal with robot swarms. Thus, we need to educate the next generation of scientists by offering the ability to simulate a multiple and diverse sets of robot swarms. Simulation mitigates operational risk prior to real-world deployment even as it allows researchers to interact with swarms while conducting actions and observing the swarm's responses to its environment. Furthermore, simulation on a cluster of remotely-located computers allows users to conduct such research in a cost-effective manner on platforms with computational capabilities far in excess of those provided by a single workstation. The question we ask, would it be possible to simulate tens of thousands of robots in a cloud environment?

However, deployment of an interconnected swarm of virtual robots in a cloud imposes additional requirements on the scientist researching such swarms. Although researchers can choose from a multitude of cloud providers currently, provisioning and configuring multiple computers presents a time and resource challenge versus a single-host setup. In addition, network se-

curity measures, such as ssh keys and port access, impede ROS' intra-cluster communication capabilities. Since these aspects typically remain out of reach for entry-level educational communities such as college freshmen or non-technical STEM students, our project attempts to lower the barrier of managing such a *virtual swarm* to a level accessible by individual researchers as well as freshmen and STEM students. In order to address the unique requirements of a networked, remotely-located swarm, we develop a multi-platform system to automate the creation and deployment of the virtual swarm, with access to numerous clouds enabled by *cloudmesh* [?].

2. VIRTUAL ROBOT SWARM COMPONENTS

2.1. Robot Operating System (ROS) [1]

The Open Source Robotics Foundation's middleware product *Robot Operating System*, or ROS, provides a framework for writing operating systems for robots. ROS offers "a collection of tools, libraries, and conventions [meant to] simplify the task of creating complex and robust robot behavior across a wide variety of robotic platforms" [3]. The Open Source Robotics

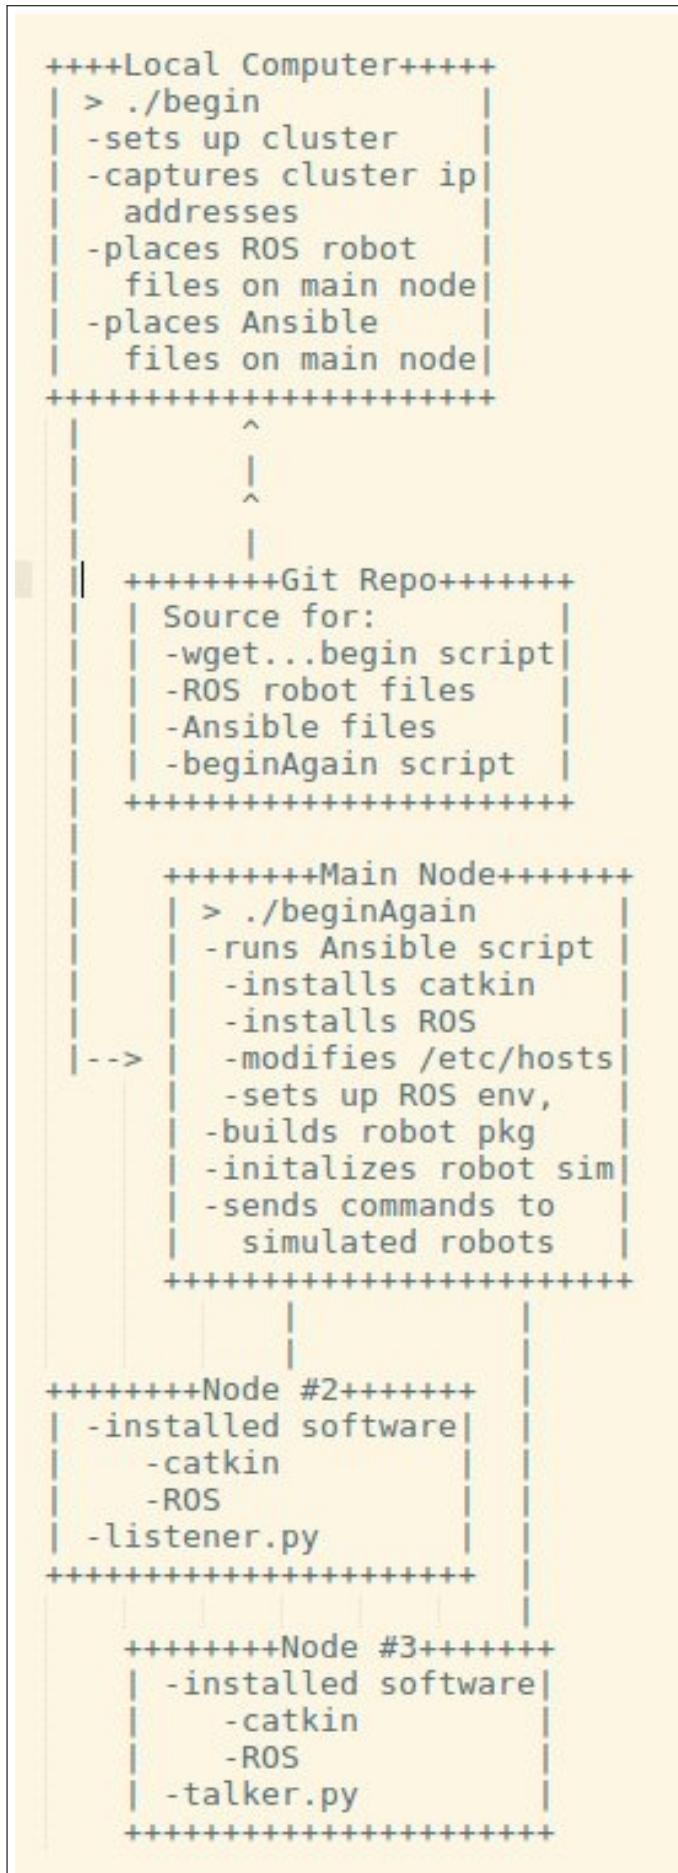


Fig. 2. Deployment Workflow for the Virtual Robot Swarm Project

programs [8].” Gazebo’s usefulness center on three main features: a) physics engines compatibility; b) its graphics engine; and c) its sensor-data generators. with respect to physics engine compatibility Gazebo interfaces well with *Open Dynamics Engine* [9] (ODE), the default; b) *Bullet* [10]; *SimBody* [11]; and, *DART* [12]. Roboticians also benefit from its 3D graphics engine, *Object-oriented Graphics Rendering Engine* [13] (OGRE), which provides a C++ class library to “[abstract away] the details of using the underlying [graphics] system libraries like Direct3D and OpenGL [14].” Finally, Gazebo can supply *sensor* data to the virtual robot. Virtual sensor support ranges from 2D cameras to Kinect-style sensors. The system can also generate *noisy* data to better simulate real-world results.

Gazebo exists as a stand-alone project, suitable for use by programs other than ROS. However, it integrates tightly with ROS given its common ownership. In fact, the version supplied with a ROS installation automatically establishes communications between Gazebo and ROS for the end-user [15].

2.3. Ansible

Red Hat, Inc.'s [16] *Ansible* software purports to simplify numerous information technology tasks. It claims to do so by a) relying upon a human-readable script syntax, YAML; and b) by automating definable and repeatable IT tasks, such as configuration management and application deployment. *Ansible*'s developers adopted a theater metaphor to describe the program's core functions. Thus, a computer's main duty within an IT infrastructure corresponds to the *role* an actor or actress might play in a theatrical production. *Ansible* calls the script a *playbook*, while the lines and directions within the script are referred to as *tasks*. Other aspects diverge from the metaphor, such as group vars and the config file (`ansible.cfg`). However, the *inventory* file hews to the metaphor - it represents the cast billing, the delineation of who plays what role. When used with an *Ansible* playbook, the inventory file specifies which servers belong to which logical group(s), i.e., which role(s).

As a result the software’s applicability extends well beyond simplistic tasks even though Ansible’s designers strive for simplification. In fact, an Ansible user can exercise fine-grained control over nearly every aspect of his or her IT infrastructure with a well-designed playbook.

Ansible also attempts to ease the burden of the IT administrator by eschewing SSL signing servers, daemons or client software. It simply pushes small programs to the target computers through an SSH connection to execute the desired tasks. When the task completes, Ansible removes the programs.

2.4. cloudmesh client toolkit

The *cloudmesh client toolkit* (cm) attempts to abstract away the complexities of establishing and utilizing different remotely-accessed computers and computer clusters [17]. Users can create, access and destroy a virtual machine or cluster of machines by issuing a single line of commands from a terminal emulator. cm supports access to clouds based on various back end-software stacks, including SLURM, SSH, Openstack and Heat. It provides an API, a command line client and a shell client.

2.5. Testing Environment

The Chameleon project's cluster of 650+ multi-core computer nodes, a joint venture between the University of Chicago and the Texas Advanced Computing Center, provides the infrastructure for the project. It is colloquially referred to as Chameleon Cloud

or CC. A 100Gbps connection runs between the two centers. Project development primarily occurred on three-node clusters created as needed using Chameleon Cloud's *m1.medium* flavor of Ubuntu 16.04. Deployment testing also utilized the *m1.small*, *m1.xlarge* and *m1.xlarge* flavors CC provides.

Table 1. Deployment Resources

Chameleon Cloud Flavor Specifications			
Flavor	vCPU	RAM	Hard Drive
m1.small	1	2	20
m1.medium	2	4	40
m1.large	4	8	80
m1.xlarge	8	16	160

vCPU: count of virtual CPUs; RAM units: MB; Hard Drive units: GB

2.6. Robots and Worlds with ROS

A ROS robot may be extremely simple, such as the pre-built talker and listener robots supplied in the ROS distribution, or as complex as the roboticist desires. If users desire to visualize the robot, s/he must define his / her ROS robot in *Universal Robot Description Format*. A ROS robot's specifications reside in a series of files with *xacro* extensions, referred to as URDF files. These files, which borrow XML's syntax structure, can include specifications for the basic shapes of the robot and its appendages (if applicable), e.g., rectangular box, cylinder, etc. with attached wheels; various kinds of joints, such as ones that rotate around an axis or extend along an axis; types and numbers of sensors; number of joints in an appendage; colors of the robot(s); etc., etc. In addition, any file in the robot description stack may reference an external file to complete the description of the robot. This capability allows the user to re-use code by incorporating previously-completed robot features into a new robot. Finally, simulations include a world file to provide the setting, or environment, in which the robot simulation occurs. As with the URDF files, the world file can be simple, like the *empty_world* file included in the ROS distribution or extremely complex.

Finally, a simulation package will include one or more *launch* files. These files coordinate the disparate aspects of the entire simulation package in a manner similar to that of Ansible's playbook file. According to the OSRF, best practices for robot simulation dictate that the main launch file consist of little more than calls to other launch files in the package [18]. Continuing with the playbook analogy, the launch files called from the main launch file would correspond to the roles defined in Ansible. This functionality significantly increases ROS' usability since managing even relatively straightforward robots, simulated or real, can prove unwieldy. For instance, this project, a proof-of-concept project, utilizes a *.world* file, four URDF files and four launch files in addition to the pre-compiled talker and listener robots.

This project uses two straightforward differential drive robots sourced from an online tutorial entitled *Simulating Robot Models in ROS (part 1)* [19]




Fig. 3. Example of a Complex Simulated Robot




Fig. 4. Example of Two Simpler Robots

3. VIRTUAL ROBOT SWARM PROJECT IMPLEMENTATION

3.1. VR Swarm task

Although the swarm accomplishes the seemingly-trivial task of driving in circles, the real accomplishment rests in proving the feasibility of automating the deployment of multiple controllable virtual robots on a remote cluster of computers and then visualizing them on a local computer.

3.2. Deployment

Achieving the aforementioned task requires coordinating a modestly-complex mix of shell commands, cloudmesh commands, Ansible commands and ROS / Gazebo commands. A shell script provides the automation for the shell, cloudmesh and ROS commands, while Ansible's *playbook* functionality handles the Ansible-focused portion. Only the initial step, retrieving and sourcing the shell script that begins the deployment process, requires user intervention. However, if the user wants to listen to the talker robot (*talker.py* and *listener.py*), s/he needs to follow the steps described in the README file.

wget...begin - Retrieve the Initialization Bash Script Deployment begins when the user retrieves the initialization file, named *begin*, from the project's Github repository. The user then starts the execution sequence by sourcing the file with the command

```
> ./begin
```

from the directory where the *begin* bash script resides.

begin bash script The bash script calls three cloudmesh_client commands:

1. > cm cluster define -n rosA1 -c 3
2. > cm cluster allocate
3. > cm cluster cross_ssh

The script creates and prepares a three-node cluster named *rosA1*. We then use two additional Cloudmesh commands:

1. > cm cluster nodes
2. > cm vm ip show

to capture the public and private ip addresses of the cluster, as well as the host names.

The ip addresses and names are written to separate files for distinct purposes. The host names and private ip addresses are used to create a file to append to each cluster node's *known_hosts* file. The private ip addresses also end up in a different file that serves as the basis for a custom Ansible inventory file.

The private ip addresses, termed static ips by Chameleon, must be used because ROS nodes communicate by binding to any available port. That is, ROS does not specify the port beforehand, and it does not consider whether or not the firewall blocks the port for security purposes. As a result, intra-cluster communication requires access to every port, i.e., the firewall cannot close any port on any computer running ROS when that computer needs to communicate with another computer in the cluster. Obviously, this presents a major security concern, especially on shared resources. Refer to the *Deployment Obstacles* section below for the workaround used with this project.

As its penultimate step, the *begin* script places all the necessary files for the demonstration on the correct cluster nodes. The ROS robot files and the Ansible files go onto the main node, while the script places a copy of the *known_hosts* addendum onto each cluster node. In addition, the script uses the *wget*

command to copy the bash script for the next step in the process onto the main node.

The final lines of the file initiate the next step in the process, running the *beginAgain* script on the master node, and perform a few administrative duties.

beginAgain bash script This script handles the software installation and initialization of the robots. It first establishes the veracity of the main node and the other nodes by connecting to each one in turn via ssh connection. If these steps do not occur before Ansible runs, the software installations never occur because Ansible hangs mid-process. Ansible installs all the necessary software on each cluster node. It installs Linux's *tree* package because the author prefers to use it to view directory structures; the ROS package; the *roshash* package; the *rosinstall* package; and, the *catkin_tools* package. Ansible also creates a new *known_hosts* file for each computer node.

Upon completing the installation, the file's instructions initialize the robot workspace. The *catkin_tools* package accomplishes this on behalf of ROS. *catkin_tools* represents an evolution of ROS' built-in *catkin* family of commands. Initializing the robot workspace essentially involves the addition of numerous hidden helper files, which assist in building the robot package. Creating the robot package, referred to as *building* the package, involves a series of automated *CMake* commands.

With the robot package built, the script copies another script from github. This small script starts ROS' *talker* robot, a simple robot included as part of ROS' tutorial packages. It relies upon Linux's *byobu* program to set up usable working environment if the user chooses to start the listener robot on the third cluster node.

Finally, execution of the last few lines of the script occurs. These lines start ROS, start the two robots, create a simulated world for the robots and then issue a command for the robots to drive in circles. Gazebo's simulation of the robots and their world come back through the ssh connection to the end user so s/he can see the simulation in real time.

talkListen bash script The *talkListen* script sets up a multiplexed terminal environment on the second node and starts ROS' *talker* robot, which ROS' maintainers include as a pre-built package with the distribution. It relies upon Linux's *byobu* terminal multiplexing program. Using a series of *> byobu -tmux ...* commands, it creates a terminal screen with three panels. ROS's *rosmaster* program starts in panel 1, while the *talker* program/robot/node starts in Panel 2. The final panel, panel 0, connects to the third node in the cluster in anticipation of the user starting the listener program (also included in the ROS distribution).

3.3. Deployment Performance

In general, completing the deployment process took around 500 seconds, or eight minutes and 20 seconds. Much of that time, around five minutes, consisted of installing ROS on the cluster nodes. Since Ansible completes each individual step of a *playbook* on each cluster node simultaneously, users should be able to provision a 5-, 10- or 100-node cluster as quickly as the three-node cluster used in this project, barring network communication bottlenecks, etc.

The total amount of time needed increased when using a higher-spec Chameleon Cloud flavor, as shown in the nearby table. However, the first portion of the process, handled by the *begin* script consistently takes about 130 seconds, irrespective of the node's vCPU/RAM/HD combination.

```

cc@eunosm3-367 (192.168.0.214) - byobu
File Edit Tabs Help

* Management: https://landscape.canonical.com
* Support: https://ubuntu.com/advantage

Get cloud support with Ubuntu Advantage Cloud Guest:
http://www.ubuntu.com/business/services/cloud

53 packages can be updated.
27 updates are security updates.

Last login: Sat May 6 04:41:16 2017 from 192.168.0.215
export ROS_MASTER_URI=http://192.168.0.214:11311
cc@eunosm3-368:~$ export ROS_MASTER_URI=http://192.168.0.214:11311
cc@eunosm3-368:~$

PARAMETERS
* /roscd: /opt/ros/kinetic
* /rosversion: 1.12.7

NODES
auto-starting new master
process[master]: started with pid [28662]
ROS_MASTER_URI=http://eunosm3-367:11311/

setting /run_id to a709fb18-3216-11e7-ac43-fa163ec
96acc
process[roscout-1]: started with pid [28751]
started core service [/roscout]

[INFO] [1494046003.522244]: hello world 1494046003.
52
[INFO] [1494046003.622244]: hello world 1494046003.
62
[INFO] [1494046003.722242]: hello world 1494046003.
72
[INFO] [1494046003.822246]: hello world 1494046003.
82
[INFO] [1494046003.922193]: hello world 1494046003.
92
[INFO] [1494046004.022249]: hello world 1494046004.
02
[INFO] [1494046004.122244]: hello world 1494046004.
12

[16.04] 0:ssh* 53!! 85m 0.01 4x2.36Hz 7.867% 2017-05-06 04:46:43

```

Fig. 5. The Multiplexed Talk-Listen Terminal

Furthermore, deployment failures seem to occur about 50% of the time, independent of the CC flavor chosen. Anecdotally, then, vCPU/RAM/HD combinations do not seem to be the first suspect to investigate with respect to failed deployments.

Table 2. Virtual Swarm Deployment Times

Time Needed to Deploy (mean; seconds)				
Flavor	Part 1	Part 2	Total	Notes
m1.small	144	441	555	4/6 attempts succeeded
m1.medium	129	358	487	2/6 attempts succeeded
m1.large	134	401	534	3/6 attempts succeeded
m1.xlarge	—	—	—	no successes

Part 1: Consists of *begin* shell script

Part 2: Consists of *beginAgain* shell script, which includes the Ansible playbook and installation of ROS

3.4. Deployment Obstacles

Intra-Node Communication In order to maintain security and enable intra-cluster communication, two of the three computer nodes must have their public ip addresses removed, i.e., disassociated, using Chameleon's browser-based graphical user interface (GUI), Horizon. Then, and only then, the ports can be opened by applying the *ros* security group to the two nodes in question (the private nodes). Since the third node, the main node, still has a public-facing ip address, the other two nodes can be accessed via an ssh connection (*ssh'ing*) to the main node and then *ssh'ing* to one of the private nodes.

Deployment Problems Even if a script or a user completes each step correctly to set up a controllable virtual swarm on remotely-located resources, problems still occur that cause deployment

failures. Usually, the issues can be resolved with manual intervention, so the deployment can be completed.

Adding the ROS Repository Key The deployment often went awry during the Ansible script step that retrieves the repository key for the ROS repo. When the key retrieval would fail, the ROS software installation would also fail. Obviously, if the main node lacks ROS, the simulation will not launch either. As a workaround, the script adds the repository to each machine with the *deb [trusted=yes] ...* syntax. Using this method obviates the need for the repo key. Although it works, applying it to repos in general creates unnecessary security risks. Therefore, the workaround should be used sparingly.

Starting the Robot Simulation Errors while starting the simulation occurred much more frequently during the development process than the key repo problem. In fact, approximately 50% of tested deployments failed, almost always at the final step when launching the simulation.

ROS prints error messages in red to the console, so they stand out. The most common error: *[gazebo_gui-3] process has died*. Sometimes, an almost identical error occurs, *[gazebo-2] process has died*, while other times the launch program cannot find parts of the robot model. The latter error consists of one or both of the following: *Error [parser.cc:581] Unable to find uri[model://ground_plane and / or Error [parser.cc:581] Unable to find uri[model://sun*. Unlike the previously-described error, this one requires manual intervention. The user needs to enter the following commands at the terminal prompt:

1. > killall roscore
2. > killall rosmaster
3. > killall gzclient
4. > killall gzserver
5. > **source** ~/.bashrc
6. > ./beginAgain

Somewhat maddeningly, this process often needs to be completed twice before the simulation will start successfully. If the simulation still will not start, refer to the instructions in the *Initializing the Swarm Manually* section.

3.5. Initializing the Swarm Manually

If needed the user can manually start the simulation with the commands listed below. You will need three terminals for the three commands, so ssh into the main node three times from three separate consoles or use *byobu* to split the terminal into three panes. The command

```
> roslaunch mybot_gazebo mybot_world.launch
```

starts the simulation. Entering

```
rostopic pub /robot1/cmd_vel geometry_msgs/Twist '{linear: {
```

```
x: 0.2, y: 0.0, z: 0.0}, angular: {x: 0.0, y: 0.0, z: 0.1}}'
```

in a free terminal will command the first robot to begin turning in a circle. Likewise,

```
rostopic pub /robot1/cmd_vel geometry_msgs/Twist '{linear: {
```

```
x: 0.2, y: 0.0, z: 0.0}, angular: {x: 0.0, y: 0.0, z: 0.1}}'
```

moves the second robot. Alternatively, one can delete the cluster and run the deployment script again.

4. VR SWARM PROJECT CONCLUSIONS

Automating the deployment of a virtual robot swarm on a cluster of remotely accessed computers should allow roboticists to more easily explore the benefits of collecting data from multiple sources at the same time over a given geographic area and timeframe. The virtual swarm presented in this paper provides framework for obtaining those benefits. It also highlights the challenges presented with an automated deployment, including a) network communication issues like obtaining the ROS repo key for each cluster node; b) compute resource security concerns since ROS nodes bind to any available port; c) the exacting nature of the interplay between the host OS, ROS and Gazebo as shown by the high number of deployment failures despite seemingly-duplicate processes and resources. The process developed for this project could perhaps benefit from the use of different tools and would undoubtedly benefit from a much-deeper understanding of how ROS functions.

The logical next step in the automation development process would seem to be projecting a differential drive robot started on one of the secondary nodes into the world created by ROS on the main node. The literature the technical feasibility of the task, although the implementation may prove more difficult than expected.

5. SUPPLEMENTAL MATERIAL

REFERENCES

- [1] Matthew Lawson and Gregor von Laszewski, "Robot Operating System (ROS)," in *Projects in Big Data Software and Applications*, G. von Laszewski, Ed., Department of Intelligent Systems Engineering, Indiana University. Indiana University, 2017, pp. 6–10.
- [2] H. Boyer, "Open Source Robotics Foundation And The Robotics Fast Track," web page, nov 2015, accessed 19-mar-2017. [Online]. Available: <https://www.osrfoundation.org/wordpress2/wp-content/uploads/2015/11/rft-boyer.pdf>
- [3] Open Source Robotics Foundation, "About ROS," Web page, mar 2017, accessed 16-mar-2017. [Online]. Available: <http://www.ros.org/about-ros/>
- [4] National Instruments, "A Layered Approach to Designing Robot Software," Web page, mar 2017, accessed 18-mar-2017. [Online]. Available: <http://www.ni.com/white-paper/13929/en/>
- [5] Open Source Robotics Foundation, "Why ROS?: Features: Core components," Web page, mar 2017, accessed 17-mar-2017. [Online]. Available: <http://www.ros.org/core-components/>
- [6] Open Source Robotics Foundation, "ROS graph concepts: Messages," Web page, aug 2016, accessed 18-mar-2017. [Online]. Available: <http://wiki.ros.org/Messages>
- [7] Open Source Robotics Foundation, "rostopic: Package Summary," Web page, jun 2016, accessed 09-apr-2017. [Online]. Available: <http://wiki.ros.org/rostopic>
- [8] Open Source Robotics Foundation, "Beginner: Overview: What is Gazebo?" Web page, apr 2017, accessed 30-apr-2017. [Online]. Available: http://gazebo-sim.org/tutorials?cat=guided_b&tut=guided_b1
- [9] Open Dynamics Engine, "Open Dynamics Engine Wiki," Web page, feb 2017, accessed 30-apr-2017. [Online]. Available: https://www.ode-wiki.org/wiki/index.php?title=Main_Page
- [10] Real-Time Physics Simulation, "Bullet Physics Library: Real-Time Physics Simulation," Web page, apr 2017, accessed 30-apr-2017. [Online]. Available: <http://bulletphysics.org/wordpress/>
- [11] P. E. Michael Sherman, "Simbody: Multibody Physics API," Web page, apr 2017, accessed 30-apr-2017. [Online]. Available: <https://simtk.org/projects/simbody/>
- [12] Georgia Tech and Carnegie Mellon University, "DART: Dynamic Animation and Robotics Toolkit," Web page, mar 2017, accessed 30-apr-2017. [Online]. Available: <http://dartsim.github.io>
- [13] Torus Knot Software Ltd, "OGRE3D," Web page, apr 2017, accessed 30-apr-2017. [Online]. Available: <http://www.ogre3d.org>
- [14] Torus Knot Software Ltd, "OGRE: About," Web page, apr 2017, accessed 30-apr-2017. [Online]. Available: <http://www.ogre3d.org/about>
- [15] Open Source Robotics Foundation, "gazebo-ros-pkgs: Package Summary," Web page, aug 2016, accessed 30-apr-2017. [Online]. Available: <http://wiki.ros.org/gazebo-ros-pkgs>
- [16] Red Hat, Inc., "redhat," Web page, apr 2017, accessed 30-04-2017. [Online]. Available: <https://www.redhat.com/en>
- [17] G. Laszewski, von, "Cloudmesh Client Toolkit," Web page, apr 2017, accessed 30-apr-2017. [Online]. Available: <http://cloudmesh.github.io/client/>
- [18] Open Source Robotics Foundation, "Roslaunch tips for large projects," Web page, nov 2012, accessed 05-may-2017. [Online]. Available: <https://goo.gl/3xBx0t>
- [19] moorerobots.com, "Simulating Robot Models with ROS (part 1)," Web page, sep 2016, accessed 15-apr-2017. [Online]. Available: <http://moorerobots.com/blog/post/1>

AUTHOR BIOGRAPHIES

Matthew Lawson received his BSBA, Finance in 1999 from the University of Tennessee, Knoxville. His research interests include data analysis, visualization and behavioral finance.

The authors would like to thank Mark McCombe for his timely and useful contribution of code snippets to the project. We greatly appreciate his time-saving additions to the cause.