Detecting Stop Signs in Images and Videos in a Robot Swarm

RAHUL RAGHATATE[1,*] AND SNEHAL CHEMBURKAR[1]

[1]*School of Informatics and Computing, Bloomington, IN 47408, U.S.A.*
Corresponding authors: rraghata@iu.edu, snehchem@iu.edu

This project is designed to deploy a scalable stop sign detection algorithm to process real-time image and video streams. The deployment is automated allowing for minimal user-interaction. Spark on Yarn provides the distributed computing power required to scale the stop sign detection algorithm to process big data. This system is useful in automated driving vehicles and advanced driver assisted systems (ADAS) to detect and classify the street signs and control the vehicle accordingly. A comparative benchmark is developed based on the performance of the application on multiple cloud systems.

Keywords: Street Signs, Ansible, Video Streams, OpenCV, Spark, Cloud

Report: https://github.com/cloudmesh/sp17-i524/tree/master/project/S17-IR-P003/report/report.pdf
Code: https://github.com/cloudmesh/cloudmesh.street

1. INTRODUCTION

There are many applications developed based on simple idea of object detection, like auto-tagging pictures (e.g. Facebook, Phototime), counting the number of people in a street (e.g. Placemeter), tracking an object in video streams, detecting vehicles to name a few. Based on this concept of object detection, we deploy a scalable software for stop sign detection using Spark on multiple clouds. The software deployment is automated using Ansible. Cloudmesh provides simple command line interface for defining the number of clusters as well as deploying them. Benchmarks are developed based on the performance of this software on different cloud systems. The database of street signs will be restricted to US street signs. The only publicly available dataset for US traffic signs is the LISA dataset [1] which is very huge. Hence, the video streams used for this project are captured by us using a mobile camera.

OpenCV is a computer vision library used to process video streams in Python. A lot of computing power goes into processing videos in real-time, this is where the cloud systems come in. We leverage the distributed computing power of Spark on Yarn for faster processing of images and videos. This solution is deployed on two different clouds to benchmark their performance. In this era of autonomous driving and advanced driver assisted systems (ADAS), detection and classification of traffic signs is a handy feature. Benchmarks have been created for the traffic sign detection on the German and Belgium Traffic Sign Datasets [2].

2. REQUIREMENT ANALYSIS

The following technologies are used throughout the project:

- Cloudmesh provides command line interface to connect and deploy clusters to different cloud systems.

- Ansible is an agentless, automated software deployment tool.

- Python - Programming language.

- Spark - Distributed computing engine.

- YARN - is the resource manager for Spark.

- OpenCV - Image and video analysis for street sign detection using open source computer vision libraries.The OpenCV library provides several transformations that can be applied to images(apply filters, transformation), detect and recognize objects in images.

3. SCOPE

The initial project idea was to automate the deployment of street sign detection algorithm over multiple cloud systems. As we proceeded through the project, we realized that training of Haar Cascade classifier is challenging. For a training dataset of 1000 samples the training can go on for 3-4 days. It turned out to be an exhaustive process. The resultant classifiers were unable to detect the specific signs. The details of the OpenCV training process we followed are given in the appendix. Due to difficulty training classifiers for street signs, we had to reduce the scope of this project to detect only stop signs using a pre-trained classifier available on Github []. The stop sign detection is implemented for both images and videos using Spark.
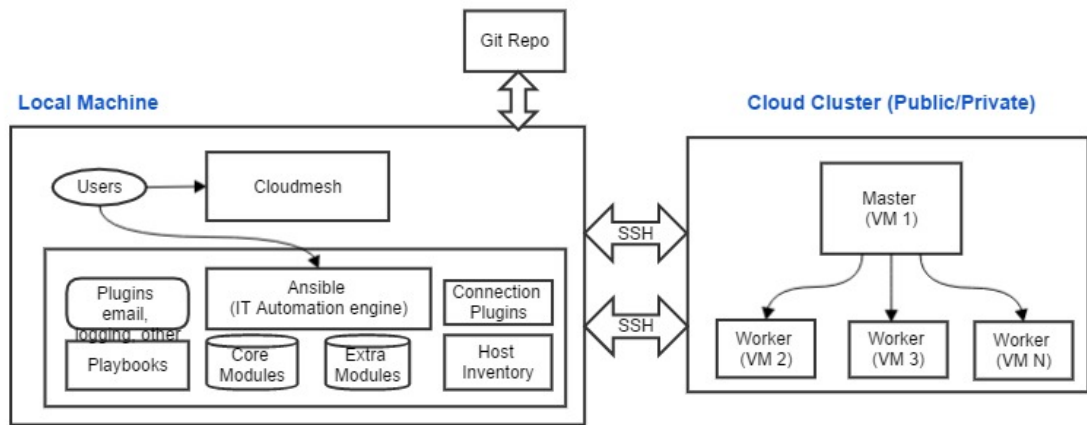
Fig. 1. System Architecture

4. SYSTEM ARCHITECTURE

Figure 1 shows an overview of our system architecture. Ansible and Cloudmesh client are installed on the host or local machine. Roles are defined in the Ansible playbook for each of the different steps in the deployment process. We execute the Ansible playbooks to instantiate the cloud machines, deploy Hadoop and Spark on them and then carryout the stop sign detection on Spark using Yarn resource manager. Once the job is submitted to Spark, the driver program initializes the SparkContext object which is responsible for the execution of the job. Input data is parallelized and sent to the worker nodes for processing. Yarn acts as the resource manager and provides executors to the worker nodes. The output is saved to the local file system on the master node and transferred to the host machine through a script. More details on the mechanics can be understood in the following sections.

5. CLOUD INFRASTRUCTURE

For the purpose of this project, we have been provided with two clouds – Chameleon [3] and Jetstream [4]. Chameleon cloud is a National Science Foundation funded experimental testbed that provides large scale cloud services to "members of the US computer Science research community and its international collaborators [3]." One can create virtual machines and manage them through the OpenStack Horizon interface. Jetstream allows researchers to leverage the computational power of cloud while retaining the look and feel of our home machines. Jetstream adds cloud based computational power to the national cyberinfrastructure [4]. Both Chameleon and Jetstream provide a cloud computing environment to researchers based on OpenStack [4]. The comparison of hardware specifications for the two cloud systems is given in table 1.

6. CLOUDMESH

Cloudmesh provides an easy interface to multiple clouds such as Chameleon and Jetstream through the command line. Cloudmesh client can be installed via pip. It is a lightweight utility that enables users to connect to different clouds from their laptops or computers. Users can customize Cloudmesh client to

	Chameleon	Jetstream
CPU	Xeon X5550	Haswell E-2680
cores	1008	7680
speed	2.3GHz	2.5GHz
RAM	5376GB	40TBr
storage	1.5PB	2 TB

Table 1. Comparison of Cloud Specification [4] [3]

suite their needs of cyberinfrastructure. It provides simple command line scripts to deploy Hadoop with Spark addon to either of the clouds mentioned in section 5. A set of cluster machine instances can be defined using command:

```
cm cluster define --count 3
```

Cloudmesh commands to create and deploy Hadoop clusters with Spark are included as tasks in the ansible playbooks to automate the deployment.

7. ANSIBLE AUTOMATION

Ansible is an easy to use, opensource automation tool that is used to automate the deployment of our project on the cloud infrastructure. Ansible is an agentless tool, that is, it does not require ansible to be deployed on the remote machines. It runs only on the host machine to deploy the required processes to the remote machines through SSH authentication. Using Ansible, we can create modules for each step of the deployment process and define the roles individually. An inventory file is used to define the machines in groups as required. A sample inventory file looks like:

```
[master]
192.128.0.1
[workernode]
192.168.0.2
192.168.0.3
```

Roles are defined in Ansible for the deployment of Hadoop-Spark cluster, environment setup on the virtual machines, stop sign detection as well as to fetch the results back to the local.

8. PYTHON-OPENCV

Python is preferred due to ease of use and familiarity over other programming languages. OpenCV also provides python library to enable object detection in python using this computer vision library. The initial scope included training a Haar Cascade Classifier to detect traffic signs and then testing the classifier on test data. But as explained in appendix section - Training a Haar Cascade Classifier - a decent classifier could not be trained. The stop sign detection algorithm utilizes a pre-trained stop sign classifier available on Github to perform the detection in images and videos. The results of the algorithm are saved as images in the `~/github/cloudmesh.street/ansible/output/output/` folder on your local machine, assuming that the git repository is cloned to `~/github/cloudmesh.street/`. The output files will have a bounding box around the detected object. The signdetectionusingspark.py file is used to process both images and videos. Based on whether the input is either image or video, the path to these files has to be modified in the spark-submit task of ansible playbook.

9. SHELL SCRIPT

Shell script is used to time each of the deployment step for benchmarking. Shell script is also useful in cases where a the deployment terminates in an error and needs to be continued from some intermediate steps. Individual shell scripts are created for each tasks in case of such issues to allow execution from the point of interruption.

10. BENCHMARK

Benchmarks are created based on the performance of the software in different cloud environments. The benchmarks for Jetstream are limited due to issues with their cloud. As the test data is images and videos, spark ran into memory issues on the m1.small flavor. Hence, we have done benchmakring using the medium and large flavours only. Tables 2 and 3 reflect that the configurations for the same flavors on the two clouds are different. We cannot directly compare the performance if the two machines have different specifications. Benchmarking is done for the deployment and the analysis on Chameleon and Jetstream.

JetStream	VCPU	RAM(GB)	Storage(GB)
m1.small	2	4	20
m1.medium	6	16	60
m1.large	10	30	60

Table 2. Jetstream Flavour Specifications

10.1. Chameleon Cloud Benchmarks

10.1.1. For flavor m1. medium and 50 Test Images

Figure 2 shows the time taken for analyzing the 50 images on Chameleon cloud for 1, 2, 3, 4, and 6 node clusters. We can see that as the number of nodes increases the time taken to analyze

Chameleon	VCPU	RAM(GB)	Storage(GB)
m1.small	1	2	20
m1.medium	2	4	40
m1.large	4	8	80

Table 3. Chameleon Flavour Specifications

the images reduces. Figure 3 reflects the total time required to complete the deployment. We can see that the total deployment time doesn't vary much upto 4 nodes but there is a steep increase in the deployment time for 6 node cluster.

Fig. 2. Time taken by sign detection task - 50 test images

Fig. 3. Time taken for complete deployment - 50 test images

10.1.2. For flavor m1.medium and a Video Input

The video input file tested on the medium flavor is just 2 sec long but since it gets converted to frames, it comes out to 53 images that are sent to spark for processing. Figure 5 shows the time taken for analyzing a single video that is 2 seconds long on Chameleon cloud for 3, 4, 6, and 7 node clusters. We can see that as the number of nodes increases the time taken to analyze the images reduces a lot. Figure 4 reflects the total time required to complete the deployment. We can see that the total deployment

time doesn't vary much upto 4 nodes but there is a steep increase in the deployment time for 6 node cluster.

10.1.3. For flavor m1.large and a Video Input

The video input file tested on large flavor is 5 sec long and after extracting the frames, it comes out to 120 images that are sent to spark for processing. Figure 7 shows the time taken for analyzing a single video that is 5 seconds long on Chameleon cloud for 1, 2, 3, and 4 node clusters. We can see that as the number of nodes increases the time taken to analyze the reduces which is expected. Figure 6 reflects the total time required to complete the deployment. We can see that the total deployment time increases steeply at first and then starts to normalize.

Fig. 6. Time taken for complete deployment - 1 Video (5 sec)

Fig. 4. Time taken for complete deployment - 1 Video (2 sec)

**Fig. 7.** Time taken for sign detection task - 1 Video (5 sec)

Fig. 5. Time taken for sign detection task - 1 Video (2 sec)

10.2. Jetstream Cloud Benchmarks

8 shows the time taken to the complete deployment on Jetstream when the number of images parsed are 4. It can been seen from the graph that as the number of nodes is increased the processing time is reduced. 9 and 10 reflect the performance of Jetstream cloud for 4 and 50 input images respectively. Sufficient data could not be gathered for Jetstream due to some issues with Jetstream. Hence we cannot conclude much about the performance of Jetstream.

Fig. 8. Total time required to deploy for 4 test images

Fig. 9. Time taken by each task for 4 test images

11. USE CASES

1. Street Sign Detection for autonomous vehicles.

2. Analysis of traffic signs in Google Street View to estimate all signs ahead hence, useful in ambulance , fire brigade services, simplest path finder etc.

Fig. 10. Time taken by each task for 50 test images

12. FUTURE WORK

This work can be expanded to detect and classify all the U.S traffic signs which can be adapted for advanced driver assisted systems (ADAS) . Moreover, efficiency of sign detection over cloud can be increased by effective distribution of data, for e.g using Hadoop distributed file system. An extension of the stop sign detection in video streams would be to output the data as videos rather than images in realtime. As the current system is scalable, benchmarks can be developed for larger dataset with multiple classifier, similar to the German and Belgium Traffic Sign Detection and Classification benchmarks [2].

13. CONCLUSION

We have been able to successfully deploy the software to Jetstream and Chameleon clouds and test their performance. On large flavors of chameleon cloud the deployment time starts to flatten out over the curve.As the number of nodes increases the time taken to deploy Hadoop and spark to those clusters increases and on the up side the analysis on Spark is faster.

ACKNOWLEDGEMENTS

REFERENCES

[1] A. Mogelmose, M. M. Trivedi, and T. B. Moeslund, "Vision-based traffic sign detection and analysis for intelligent driver assistance systems: Perspectives and survey," *IEEE Transactions on Intelligent Transportation Systems*, vol. 13, no. 4, pp. 1484–1497, Dec 2012, accessed : 2017-03-29.

[2] M. Mathias, R. Timofte, R. Benenson, and L. V. Gool, "Traffic sign recognition #x2014; how far are we from the solution?" in *The 2013 International Joint Conference on Neural Networks (IJCNN)*, Aug 2013, pp. 1–8, accessed : 2017-03-25.

[3] C. Cloud, "Chameleon," Web Page, accessed: 2017-04-19. [Online]. Available: https://www.chameleoncloud.org/docs/user-faq/

[4] I. University, "Jetstream," Web Page, accessed: 2017-04-11. [Online]. Available: https://jetstream-cloud.org/

[5] S. Soo, "Object detection using haar-cascade classifier," *Institute of Computer Science, University of Tartu*, 2014, accessed : 2017-03-29.

[6] P. Viola and M. Jones, "Rapid object detection using a boosted cascade of simple features," in *Proceedings of the 2001 IEEE Computer Society Conference on Computer Vision and Pattern Recognition. CVPR 2001*, vol. 1, 2001, pp. I–511–I–518 vol.1, accessed : 2017-03-29.

[7] T. Ball, "Train your own opencv haar classifier - coding robin," Webpage, Jul. 2013, accessed: 2017-04-05. [Online]. Available: http://coding-robin.de/2013/07/22/train-your-own-opencv-haar-classifier.html

[8] G. Bradski, "Home - opencv/opencv wiki," Code Repository, accessed: 03-12-2017. [Online]. Available: https://github.com/opencv/opencv/wiki

[9] M. Gaynor, "markgaynor/stopsigns - detecting stop signs in google street view images of a provided address." Code Repository, Sep. 2016, accessed: 2017-3-21. [Online]. Available: https://github.com/markgaynor/stopsigns

AUTHOR BIOGRAPHIES

Rahul Raghatate is a graduate student at Indiana University. He will receive his Masters degree in Data Science in 2018. His research interests include Big Data and Machine Learning.

Snehal Chemburkar is a graduate student at Indiana University. She will receive her degree in Data Science in 2018. Her research interests include Big Data and Machine Learning.

A. OPENCV IMAGE PROCESSING

OpenCV provides a range of computer vision algorithms to detect objects in images. One of the simplest method for object detection is based on color. The results of Color based detection method are largely affected by the lighting conditions and one require the user to calibrate multiple times before they might get a better result in the real world [5]. Hence this technique is not very popular when detecting objects in the real world Haar features is a sophisticated technique that uses the features specific to the object in question. It been seen that working with RGB pixel values in every single pixel in the image results in computationally expensive and slow feature calculation. "A Haar-like feature considers neighboring rectangular regions at a specific location in a detection window, sums up the pixel intensities in each region and calculates the difference between these sums. This difference is then used to categorize subsections of an image [5]." OpenCV provides a Haar feature based cascade classifier that can be used for object detection, as proposed by in [6].

A.1. Data Collection

The publicly available data set for U.S street signs is the LISA traffic dataset [1]. This dataset contains images for 47 different traffic signs. But since the data set itself is approximately 7GB, we extracted 50 images from the dataset for the purpose of testing. For our training set, we captured images of street signs and put together a few positive images for the street signs. The positive images were cropped to only contain the street sign and resized to 50x50.

A.2. Train a Haar feature-based Cascade Classifier

Based on tutorials provided in [7], [8] , we carried out multiple experiments to train a classifier to detect street signs. Since each sign needed to trained separately, we picked stop, yield, and signal ahead signs to start with. To train a classifier, we firstly required gathering at least a few positive and many negative images. The positive images are images of the object alone cropped to a size of 24x24 or 50x50 whereas the negative images should not contain the object in consideration here. In case we have a single positive image or a few positive images, OpenCV provides a utility called opencv_createsamples to generate the training and test datasets in *.vec format that is supported by the opencv_traincascade utility. The samples generated from the opencv_createsamples can be passed to the opencv_traincascade utility to get a trained classifier. Multiple experiments were carried out by differing the sample sizes (the width by height of the positive images) and varying the number of positive and negative images. As the dataset and width by height increases the computational time increases. Below are the trainings that were carried out for stop sign:

```
opencv_traincascade -data classifier -vec samples.vec
```

```
-bg negatives.txt -numStages 20 -minHitRate 0.999
-maxFalseAlarmRate 0.5 -numPos 120 -numNeg 200 -w 50
-h 50 -mode ALL -precalcValBufSize 1024
-precalcIdxBufSize 1024

opencv_traincascade -data classifier -vec samples.vec
-bg negatives.txt -numStages 20 -minHitRate 0.999
-maxFalseAlarmRate 0.5 -numPos 200 -numNeg 350 -w 50
-h 50 -mode ALL -precalcValBufSize 1024
-precalcIdxBufSize 1024

opencv_traincascade -data classifier -vec samples.vec
-bg negatives.txt -numStages 20 -minHitRate 0.999
-maxFalseAlarmRate 0.5 -numPos 600 -numNeg 100 -w 50
-h 50 -mode ALL -precalcValBufSize 1024
-precalcIdxBufSize 1024
```

When we increased the number of positive samples to 600 for the 50x50 image size, the training ran for 3.5 days. The resulting classifier was unable to detect the stop signs in the test data. After a couple more experiments and another week invested in training the classifier to no good results, we found success with a pre-trained classifier for stop signs available at [9]. The results of this classifier are shown in 11.

Fig. 11. Stop Sign Detection

After successful testing of stop sign classifier, we proceeded to train the classifiers for Yield sign and Signal Ahead sign. We trained 3 classifiers each for both these signs while increasing the number of positive images from 600, 900, 1200. Even after increasing the number of positive images up to 1200 the resulting classifiers were not efficient enough to detect the signs in images. As each training had resulted in a loss of approximately 3.5 days, we realized that this could not be covered as part of this project and restricted ourselves to the stop sign detection.

A.2.1. Outcomes

- From the many experiments we carried out, we learned that there is no fixed number of samples that will yield a decent result.

- Future work can be done on training the U.S traffic signs, since there are no classifiers available for them . With the growing market for autonomous vehicles and assisted driving technology, having trained classifiers for the traffic signs might prove to helpful.

B. EXECUTION SUMMARY

This section specifies the week by week timeline for project completion.

1. Mar 6 - Mar 12, 2017: During this week we created virtual machines on Chameleon cloud using Cloudmesh and submit the project proposal.

2. Mar 13-Mar 19, 2017: Deployed Hadoop cluster to Chameleon cloud using Cloudmesh and develop Ansible playbook to install the required software packages to the clusters (OpenCV, Python and dependencies)

3. Mar 20-Mar 26, 2017: Collated data for training and test data sets and trained stop sign classifier. Developed Ansible playbook to deploy Hadoop and Spark to the cloud machines.

4. Mar 27-Apr 02, 2017: Trained data for stop and yield sign classifiers using OpenCV. Developed Ansible playbook to setup the OpenCV python environment on Spark clusters.

5. Apr 03-Apr 09, 2017: Trained data for signal ahead sign using OpenCV. Test stop sign classifier on local machine and chameleon cloud.

6. Apr 10 - Apr 16, 2017: Tested classifier on Spark and created deployable software package using shell script.

7. Apr 17-Apr 23, 2017: Completed project report and developed benchmarks for the project.