

Cloudmesh client extension for AWS

GREGOR VON LASZEWSKI¹, MILIND SURYAWANSHI¹, AND PIYUSH RAI¹

¹School of Informatics and Computing, Bloomington, IN 47408, U.S.A.

project-006, May 11, 2017

Cloudmesh client provides an easy abstraction to gain access to hybrid clouds from the commandline. In contrast to other systems it provides a mechanism of adding plugins to the command. The command can not only be used as a commandline tool, but as a shell. Due to the swift developments in the area of cyber infrastructure such a convenient plugin allows us to integrate easily new services easily. This work uses our next generation cloudmesh shell and demonstrates the ease of integrating Amazon Web Services services into it. Hence Cloudmesh client can via other plugins achieve interoperability for virtual machine management by providing plugins for OpenStack, Azure, Comet Cloud Services and other frameworks including containers. Here we focus on extensions to bring its capability to work with Amazon EC2 services. As our integration demonstrates how to leverage libcloud as a backend, more than 30 other providers could easily be targeted.

© 2017 <https://creativecommons.org/licenses/>. The authors verify that the text is not plagiarized.

Keywords: Cloud, I524

Report: <https://github.com/cloudmesh/sp17-i524/blob/master/project/S17-IR-P006/report/report.pdf>

Code: <https://github.com/cloudmesh/cloudmesh.aws>

1. INTRODUCTION

Gregor: will include section about cloudmesh [1]

Amazon provides a web service in a cloud (AWS) to deploy virtual machines (VM) with numerous operating systems environment available in different flavors. The images are called Amazon Machine Image (AMI) and are available as pre-configured for different applications. One can also create his own custom environment [2].

The Amazon EC2 drivers provides numerous functionalities to authenticate into the AWS, list available configurations and create and boot VMs. Many interfaces exist providing access to many different frameworks and platforms. Being able to utilize access through Web services allows developers to integrate AWS into their service offerings. Within Indiana University we have provided for the last seven years cloud services based on a variety of cloud services. While observing the evolution and practical use of these services it has become clear that in order to switch easily from one cloud to another we need a uniform interfaces that provides the most elementary service offerings of managing virtual machines. Our design offers the ability to switch easily between different cloud providers. Figure ?? demonstrates this convenient feature. Here we define in our cloudmesh client a cloud and set it to AWS, we then start a vm, we can do the same for OpenStack as demonstrated. Important to note is that all details of name assignment, flavor and image

management, are abstracted out, but can be controlled and are provided as part of an easy to manage configuration framework. This has advantages for university settings as we can distribute such configurations for particular classes and customize the use for many hundreds of non cloud experts. In fact using our abstractions allows non cloud experts easily to use clouds.

```
set default cloud=aws
vm boot
set default cloud=openstack
vm boot
```

Fig. 1. cms aws image refresh

2. DESIGN

2.1. Requirements

Requirements are missing

3. IMPLEMENTATION CONSIDERATIONS

The AWS client is based on cloudmesh.cmd5 [3] and cloudmesh.common [4]. It uses mongodb in the back-end to store the cloud information such as list of available images or instances running on the cloud. Requests library [5] is used to

connect to the back-end through rest services. The rest service is deployed using cloudmesh REST framework [6].

Design Diagram is missing

```
client -> rest (eve) -> database (nosql)
      ^ object definitions
```

4. IMPLEMENTATION

Some details about REST service and auto creation from examples. Gregor may have text.

5. ARCHITECTURE

The commands are implemented as methods of a class `AwsCommand` which is based on `PluginCommand` from `cloudmesh shell`. Depending on the arguments passed, corresponding routine is called from `Aws client API` which acts as a wrapper around the `libcloud EC2 drivers` [7] and is also responsible to connect with back-end database apart from reading the configuration from `yaml` file. The entire code is written in `Python`.

5.1. Technologies Used

To adhere to the principal of code reuse we are using a number of core technologies that allow us to keep the overall developed new code small. This includes leveraging our previous *cloudmesh* efforts, the reuse of *mongodb* and the automatic creation of a REST service based on example objects with *eveengine*. Particularly the following software components have been leveraged:

Libcloud EC2 driver: The driver provides a number of functions for various functionalities such as listing the available nodes, generating a key pair and deploying a VM.

MongoDb: MongoDB is an open source document store database. It's used by AWS client to store information about various VM configuration options available on the Amazon EC2 cloud. It's also used to store information regarding the VMs that are running on the cloud. The information in the database is refreshed whenever it's fetched from the cloud.

Cloudmesh.common: A library of many commonly useful functions and classes shared among all cloudmesh projects [4].

Cloudmesh.cmd5: The `cmd5` is a dynamically extensible CMD based command shell [3].

Cloudmesh.evegenie: The schema for the collections are specified in `json` files which are then converted to `Eve` syntax using *evegenie*. It creates the configuration file for starting up the rest services using *cloudmesh rest* framework.

Cloudmesh.rest: The *cloudmesh rest* framework is used to deploy and start the *mongodb* and *rest* services. The schema for the objects to be stored in the database is collectively specified in `json` file called *all.json*. The schema defined in the file is closely associated with the code in *awsclient.py* which is responsible for fetching the information from cloud and passing it to *mongodb* through *rest* services. From our experience during the development of this project, we observed that *rest* services required the schema to be precise and didn't handle null values for collection fields. There's another file *all.settings.py* which contains the configuration information for *rest* services such as port *mongodb*

is running on along with the database name. It contains the schema for the collections and the list of methods to be provided by the *rest* services. The database connectivity was initially developed using *pymongo* library. However, during the review it was suggested that the code based on *pymongo* is quite low level and does not have adequate security features. This led to the use of *Pythhon Requests* library.

6. USAGE

One needs to create an `AWS` account first to be able to access its cloud. The instructions for it are available at [8]. A pair of access key and secret keys are required to be generated to authenticate into the cloud [9]. These keys are required to be kept confidential. They are to be specified in a `yaml` configuration file. Default `vm` image and flavor can also be specified in the configuration file (see Figure 2).

```
cloudmesh:
  clouds:
    aws:
      credentials:
        EC2_ACCESS_KEY: 'ACCESS KEY'
        EC2_SECRET_KEY: 'SECRET KEY'
        . . .
      default:
        image: 'IMAGE ID'
        size: 'SIZE ID'
        location: 'LOCATION'
```

Fig. 2. Configuration

The user need to ensure that the combination of image, size and location is valid and that it's account has the required privileges for it. The instructions for installation of `AWS` client can be found at [10]. Once, the client has been installed and configurations settings enabled, the *mongodb* and the *rest* services to access the database can be started by following command from the root directory of the code:

```
make rest
```

The above services will be required by `AWS` client to store cloud related information locally. The user can now execute the client commands e.g.:

```
cms aws flavor refresh
```

This will fetch the list of image sizes available on Amazon EC2 cloud.

7. AWS COMMANDS

8. BENCHMARK

Benchmarks missing

9. USE CASE: CREATE AN AWS NODE

The following are the pre-requisites to be able to create an EC2 node.

9.1. Prerequisite

URLS IN REFERENCES

1. User should have a valid AWS account.
2. [Create a IAM](#) user for which the access and secret key will have to be generated [11]. The keys will have to be added to the yaml cloudmesh configuration file.
3. The required permissions are needed to be granted to the user.
 - (a) visit [IAM home](#)
 - (b) select **policies** in left hand menu
 - (c) create **administrator policy** from amazon's existing policies
 - (d) select **administrator checkbox** and **attach** to your user
4. Open the `../cloudmesh.yaml` configuration file and update with `EC2_ACCESS_KEY`, `EC2_SECRET_KEY`, you just copied. Now set default flavor, image and location, as `t2.micro`, `ami-0183d861` and `us-west-1` respectively in same file.
5. MongoDB server should be up and running (please refer section 2. Getting started)

9.2. Create Node

- (a) Create keypair name using command

```
cms aws keypair create AWS1
```

It will create the *keypair name*, it is essential to create the *node*. To verify it, `cms aws keypair list` command will list down all created *keypairs* so far.

- (b) Now create the node

```
cms aws vm boot ami-0183d861
```

Above command will create a node instance with the image of `ami-0183d861`.

AUTHOR BIOGRAPHIES

Milind Suryawanshi received his BE (Electronics and Telecommunication) in 2010 from The University of Pune. His research interests include Big Data analytics for intelligence and research. **Piyush Rai** received his BE (Computer) in 2011 from The University of Pune. His research interests also include Big Data analytics for military intelligence and financial markets. **Gregor von Laszewski** has written more than 100 publications in the area of Grid and Cloud computing. He obtained a Ph.D. from Syracuse University and has worked at Argonne National Laboratory. His team is providing with cloudmesh a client interface to XSEDE SDSC comet virtual clusters.

REFERENCES

- [1] Web Page. [Online]. Available: <https://github.com/cloudmesh/client>
- [2] Web Page. [Online]. Available: <https://aws.amazon.com/ec2/details/>
- [3] G. von Laszewski, "cloudmesh.cmd5," Web Page. [Online]. Available: <https://github.com/cloudmesh/cloudmesh.cmd5>
- [4] G. von Laszewski, "cloudmesh.common," Web Page. [Online]. Available: <https://github.com/cloudmesh/cloudmesh.common>
- [5] Web Page. [Online]. Available: <https://pypi.python.org/pypi/requests>
- [6] Web Page. [Online]. Available: <https://github.com/cloudmesh/cloudmesh.rest>
- [7] Web Page. [Online]. Available: <http://libcloud.readthedocs.io/en/latest/compute/drivers/ec2.html>
- [8] Web Page. [Online]. Available: <http://docs.aws.amazon.com/AmazonSimpleDB/latest/DeveloperGuide/AboutAWSAccounts.html>
- [9] Web Page. [Online]. Available: <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-key-pairs.html>
- [10] Web Page. [Online]. Available: <https://github.com/cloudmesh/cloudmesh.aws>
- [11] Web Page. [Online]. Available: <http://stackoverflow.com/questions/28222445/aws-cli-client-unauthorizedoperation-even-when-keys-are-set>

id	name	driver
aki-02b79b47	pv-grub-hd00_1.03-1386	<libcloud.compute.drivers.ec2.EC2NodeDriver object at 0x104f29150>
aki-033c6d46		<libcloud.compute.drivers.ec2.EC2NodeDriver object at 0x104f29150>
aki-037a5e46	ubuntu/kernels-testing/ubuntu-lucid-1386-linux-image-2.6.32-347-ec2-v-2.6.32-347.52-kernel	<libcloud.compute.drivers.ec2.EC2NodeDriver object at 0x104f29150>
aki-04f3a241	ubuntu-kernels/ubuntu-lucid-amd64-linux-image-2.6.32-308-ec2-v-2.6.32-308.15-kernel	<libcloud.compute.drivers.ec2.EC2NodeDriver object at 0x104f29150>
aki-052d7c40		<libcloud.compute.drivers.ec2.EC2NodeDriver object at 0x104f29150>

Fig. 3. aws image refresh

id	name	ram	disk	bandwidth	price
t1.micro	Micro Instance	627	15		0.025
m1.small	Small Instance	1740	160		0.047
m1.medium	Medium Instance	3840	410		0.095
m1.large	Large Instance	7680	840		0.19
m1.xlarge	Extra Large Instance	15360	1680		0.370

Fig. 4. aws flavor refresh

uuid	name	state	public_ips	private_ips	provider
8f3eca1839b3fa8ed3c56f97b637f96af39c886b	test1	pending	[]	['172.31.7.176']	Amazon EC2

Fig. 5. aws vm boot ami-0183d861

uuid	name	state	public_ips	private_ips	provider
8f3eca1839b3fa8ed3c56f97b637f96af39c886b	test1	running	['54.153.119.61']	['172.31.7.176']	Amazon EC2
27a286025aa58512a91f55a1e28432fc9883f8cc	test1	running	['52.53.153.142']	['172.31.5.131']	Amazon EC2

Fig. 6. aws vm list

name	fingerprint	driver
AWSCLI	ed:12:c4:34:4a:08:97:31:09:48:37:37:ed:e a:23:5d:e5:fc:88:00	Amazon EC2

Fig. 7. aws keypair create AWSCLI

id	name	country	availability_zone	zone_state	region_name	provider
0	us-west-1a	USA	us-west-1a	available	us-west-1	Amazon EC2
1	us-west-1b	USA	us-west-1b	available	us-west-1	Amazon EC2

Fig. 8. aws location refresh

id	size	driver
vol-0a4788fb2d0a67c7e	1	Amazon EC2

Fig. 9. aws volume create VOL_TEST_1

id	size	driver
vol-0c30e77f6ca2a263a	1	Amazon EC2
vol-009775889a2d50154	1	Amazon EC2
vol-0ec3dbd48a667ab8e	1	Amazon EC2
vol-0c5f400af50f5b460	1	Amazon EC2

Fig. 10. aws volume refresh

Table 1. Cloudmesh AWS Commands

Function	Description	Command
Refresh	Whether to always fetch the information from the cloud over the network or display it from the local database when asked can be configured by setting the configuration variable <i>refresh</i> to either <i>on</i> or <i>off</i> . When the value is set to <i>on</i> , the information will always be fetched from the cloud. This functionality is implemented for only some of the commands as of now.	<code>aws refresh on</code>
Image refresh	The images available on the cloud could be listed using this command. The database will also be updated with the newly fetched list. The command may take up to minutes to run as the list contains more than 24,000 images.	<code>aws image refresh</code>
Image list	Depending on whether the value of <i>refresh</i> is set to <i>on</i> or <i>off</i> , the list is either fetched from the cloud or from the local database.	<code>aws image list</code>
Flavor refresh	This will list the different sizes of VMs that are available on the cloud. The information is fetched from the cloud and stored locally.	<code>aws flavor refresh</code>
Flavor list	The flavor list is either fetched from the cloud or from the local database depending on whether the value of <i>refresh</i> is set to <i>on</i> or <i>off</i> . <code>aws flavor list</code>	
Start/Boot vm	It creates a new node instance and start that node automatically. The required parameter is <i>IMAGE_ID</i> . This command assumes user has created keypair with name <i>AWS1</i> . The default values; flavor and location are taken from <i>cloudmesh.yaml</i> . Those values are required to be set before executing this command. Once it has been created, it can be seen using <i>vm list</i> .	<code>aws vm boot IMAGE_ID</code>
Reboot vm	This command is used to reboot a running vm instance. The state of the node gets changed from <i>RUNNING</i> to <i>REBOOTING</i> while its restart. The rebooting time of node depends on the configuration the configuration that was selected while its creation. <i>vm reboot</i> requires <i>NODE_UUID</i> of the node to be rebooted which can be retrieved using either <i>vm list</i> or <i>vm refresh</i> .	<code>aws vm reboot NODE_UUID</code>
VM delete	It destroys the instance of a node and all the data associated with it including backup. It takes few seconds to terminated the instance. At first, the state of the node changes to <i>TERMINATED</i> which can be observed using <i>vm refresh</i> . The node gets removed from the list subsequently.	<code>aws vm delete UUID</code>
VM list	It lists out all the nodes along with their status as stored in the database.	<code>aws vm list</code>
Keypair create	To create the node, one of the essential component is <i>key pair</i> . Amazon EC2 uses public key of the user to encrypt the password and then recipient uses private key to decrypt it. These public and private keys are known as a <i>key pair</i> . They are required to be given a name.	<code>aws keypair create NAME</code>
Keypair delete	It allows user to delete a created <i>key pair</i> which is no longer in use.	<code>aws keypair delete NAME</code>
Keypair refresh	It refreshes the list of created <i>key pairs</i> in the database and also displays it on the screen.	<code>aws keypair refresh</code>
Keypairs list	It lists out all the <i>key pairs</i> information stored locally in the database.	<code>aws keypair list</code>
Keypair get	It returns the key pair object, which has the name of <i>key pairs</i> , driver and hash key.	<code>aws keypair get NAME</code>

Table 2. Cloudmesh AWS Commands

Locations refresh	It will show all the available locations associated with Amazon EC2 account. For free tier, user will get two locations. More locations will be available in paid service. The database records are also updated with this command.	<code>aws location refresh</code>
Locations list	It lists the locations information stored in the database.	<code>aws location list</code>
Volume create	It creates the volume for vm. The size of volume is specified in GB, the default value is set to 1 GB. The maximum number of volumes that can be attached to a vm will depend on its operating system.	<code>aws volume create VOLUME_NAME</code>
Volume re-fresh	This command shows the created volumes with the id, size and the driver name. It also updates the database.	<code>aws volume refresh</code>
Volume list	This command displays the volume information stored in the database.	<code>aws volume list</code>
Volume delete	User can delete the unwanted volumes using the <i>VOLUME_ID</i> .	<code>aws volume delete VOLUME_ID</code>