

# Deploying a spam message detection application using R over Docker and Kubernetes

SAGAR VORA<sup>1,\*</sup> AND RAHUL SINGH<sup>1,\*\*</sup>

<sup>\*</sup> School of Informatics and Computing, Bloomington, IN 47408, U.S.A.

<sup>\*\*</sup> Corresponding authors: vorasagar7@gmail.com, rahul\_singh919@yahoo.com

Project-S17-IR-P007, November 26, 2017

Application containerization and automated deployment allows developers to focus on writing code without worrying about the environment in which their software is supposed to run. This paper aims to containerize a standalone R application and deploy it on a multi-node kubernetes cluster and benchmark application performance on different clouds. Email spam messages are a nuisance to the web community. If not filtered accurately, they may work like an email-bomb attack where a stream of spam messages shall make a user to accidentally omit a legitimate message. We use several established data mining algorithms to classify messages as legitimate or not, and compare their classification accuracy. Application and cluster performance is benchmarked and compared across Chameleon and Kilo clouds. © 2017

<https://creativecommons.org/licenses/>. The authors verify that the text is not plagiarized.

**Keywords:** Docker, Ansible, Kubernetes, R, Spam

Report: <https://github.com/cloudmesh/sp17-i524/tree/master/project/S17-IR-P007/report/report.pdf>

Code: <https://github.com/cloudmesh/cloudmesh.kubernetes>

## 1. INTRODUCTION

Email and text messaging are the most formal means of communication for internet users. Web mining helps any organization to discover identities of users that it can target to advertise their business by sending them emails, and at times without their consent. The ease with which content can be generated and published has also made it easier to create spam. Spam can be stated as any information which does not add value to a user of the web. Messages which are inappropriate, unsolicited, repeated and irrelevant can be all classified as spam. In this report, we are using data mining algorithms like SVM (Simple Vector Machine), KNN (K-Nearest Neighbor) and AdaBoost over an SMS collection dataset of 5574 messages to differentiate spam messages from legitimate ones. R [1] shall be used to develop the code as it provides supporting text mining libraries to implement these algorithms.

To automate application deployment such that it can be tested on any environment, Docker [2] is used to containerize the application along with the its dependent libraries. The cluster management tool - Kubernetes [3], is used to auto deploy the application on a multi-node cluster. Scripts are written in Ansible [4] to automate the deployment process. As Kubernetes itself scales the application container according to our specification, a deployment engineer shall only be responsible to write a yaml specification file that is read by the kubernetes engine. Application benchmarking results are achieved by running CMD5 commands that trigger ansible scripts to deploy the application

on Chameleon cloud and Kilo cloud.

## 2. ELEMENTS OF THE PROJECT

### 2.1. Ansible

Human labour has been replaced by machines and the world is moving towards an age of automation. Computer scientists prefer writing scripts to perform their daily activities instead of manually performing them again and again. Ansible is an open source automation engine that can be used to automate tasks such as cloud instance creation and deployment, configuration management, and application deployment. Apart from these basic tasks, Ansible can also be used to perform advanced tasks such as continuous deployment or rolling out updates without turning off the entire cluster.

#### 2.1.1. Architecture

Ansible doesn't require any software to be installed on the remote machine in order to control it. It implicitly manages remote machines using SSH [5] or WinRM [6], which are natively present on those platforms [7].

Ansible architecture comprises of two types of servers, the *controlling machine* and the *nodes* [8]. The controlling machine is responsible for managing the nodes over SSH. The location of these nodes is described by the controlling machine using its *inventory* [9].

### 2.1.2. Playbooks

"Playbooks are a powerful mode of using Ansible's configuration, deployment, and orchestration capabilities" [10]. Using playbooks we can govern the rules of managing and controlling remote machines. Using inventory files in conjunction with the playbooks, we can break down our ansible tasks into Ansible roles format, where each role is binded with a set of commands to be executed. The roles are defined in the playbooks which Ansible executes in a sequential fashion. Playbooks can also be used to manage configuration and deployment to remote machines.

### 2.2. Executing a Playbook

```
$ ansible-playbook xxx_playbook.yml -i inventory.txt
```

**Fig. 1.** Command to execute a playbook

The *ansible-playbook* command followed by the playbook name is used to execute the playbook. The *-i* option specifies the inventory file that is to be used in conjunction with the playbook file.

### 2.3. R

R is a language and environment for statistical computing and graphics [1]. R provides a variety of statistical model analysis, classification, clustering and graphical techniques to provide this environment. R's capability allows us to build a highly desirable analysis model for our application.

### 2.4. Docker

Docker is an open-source project that automates application deployment by packaging the application in *containers*. Containers provide application portability by bundling together an application and its needed resources in a package so that they can be deployed on different platforms without worrying about resource dependencies. Application containerization is an OS (Operating System) level *virtualization* [11] for deploying and running an application instance without launching a virtual machine for each application [12]. A container has its own environment variables, filesystem and libraries that is needed by the application, thus eliminating OS or hardware dependency. Containers abstract the OS kernel while a VM (*Virtual Machine*) [13] hypervisor abstracts an entire device.

Docker allows application developers to package their applications into isolated containers. Docker automates the repetitive tasks of setting up and configuring development environments thus allowing developers to focus only on building software. A dockerized application can simply ship between platforms as the complexity of software dependencies is handled by the container. Docker standardizes container creation and can be used to pack, ship and run an application as a lightweight container that can run in any environment. Docker can be integrated with other devOps applications like *Puppet*, *Chef*, *Vagrant*, *Ansible* and *Kubernetes*. We shall use Docker with Ansible and Kubernetes in our project.

#### 2.4.1. Dockerfile and DockerImage

To package an application and its dependencies in a single file, docker introduces the concept of a *docker image*. The docker engine creates a docker image by parsing contents of a *dockerfile*. A dockerfile is a script composed of various commands to build a container in a step-by-step, layer-by-layer manner [14]. Once an

image has been built it can be shared with other users by pushing it to a public repository on *DockerHub* or *GoogleCloudPlatform*. In this manner, an image once built by the docker engine can be used across the organization by making a docker pull request.

### 2.5. Kubernetes

Kubernetes is an open-source platform for automating deployment, management and scaling of containerized applications across a cluster [15]. The more granular an application is, the more components it consists of and thus requires organized management of these components. Kubernetes helps in faster deployment of applications and scaling them on the fly. Moreover it optimizes the use of hardware by using the resources which are needed. Kubernetes provides container management features like component replication, load-balancing, service-discovery and logging across components [16]. A Kubernetes cluster can be deployed on either physical or virtual machines. We shall be deploying a kubernetes cluster using *kubeadm* - the kubernetes command line tool that allows us to setup and maintain nodes in a cluster by parsing *yaml* structured deployment files.

### 2.6. Kubernetes Terminologies

Kubernetes defines the following set of primitives which provide mechanisms for deploying and scaling applications.

#### 2.6.1. Pods

A pod is the smallest unit of a kubernetes cluster and has a unique IP address within the cluster. A pod consists of one or more containers that can share resources and can be controlled as a single application [15] [17]. Thus, all the involved containers in a pod are scheduled on the same host. A pod can be thought of as a single virtual machine in terms of resource sharing and scheduling. Pods can be managed manually using the *Kubernetes API* or can be managed by a controller.

#### 2.6.2. Services

A kubernetes service is a collection of pods that perform the same function and are presented as a single entity. This way a service can be emphasized as a one tier of a multi-tier application. Service acts as an interface to a group of containers so that service-consumers need only reference the single access location. Kubernetes facilitates service discovery by assigning an IP address and a namespace to a service. This idea abstracts the change of IP addresses of pods within a service that result due to pod failure or pod rescheduling.

#### 2.6.3. Replication Controllers

A replication controller is a framework for horizontal scaling of pods. Semantics of a pod are defined in a *<pod\_name>.yaml* file which also defines the replication details that need to be done. The replication controller performs replication by scaling a number of pods across a cluster based on the pod definition file [15]. The replication controller has to make sure that a certain number of copies of a pod are always up and running. Thus, in event of a pod failure it replaces the failed pod with a new replica.

#### 2.6.4. Labels and Selectors

Kubernetes allows users or internal components to assign a key-value pair tag to any API object in the system. An object can have one or more labels associated with it, but each with a unique key. eg: *appversion* = 1.0, *development\_stage* = 4. Label selectors are queries against labels that return matching objects [17]. This

way each object in the system can be referenced with a single label or a combination of multiple labels for fine grained control.

## 2.7. Kubernetes architecture

Kubernetes exhibits the master-slave architecture. The components can be split into those that manage an individual node(slave) and those that manage the master or control plane.

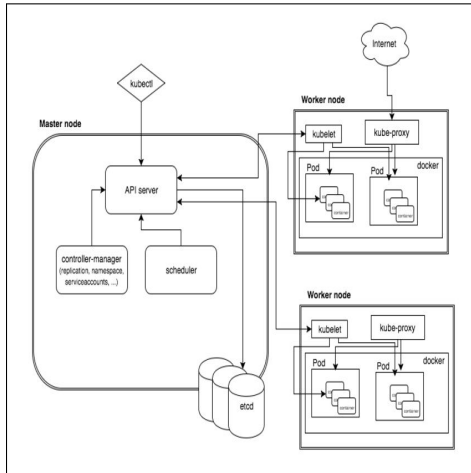


Fig. 2. Kubernetes architecture [16]

### Kubernetes Master Node/Control plane

The Kubernetes master node is responsible for managing the kubernetes cluster and orchestrating the worker nodes, where the actual pods are scheduled. The master node can also work as a single node cluster where it can schedule pods to work on the master node itself by tainting the schedule policy rule on the node. The master node also referred as the control plane consists of several components:

#### 2.7.1. API Server

The API server is the most fundamental component of the Kubernetes master and serves as the entry point for all the REST commands used to control the cluster. The API server serves up the Kubernetes API using *JSON over HTTP*, providing both internal and external interface to the cluster [17] [18]. It validates the REST requests, executes them and updates the status of the objects in the *etcd* storage.

#### 2.7.2. etcd storage

*etcd* is a simple, distributed and consistent key-value store that stores configuration data of the cluster and represents the state of the cluster at any point of time [15]. Kubernetes uses *etcd* for service discovery and provides a simple *HTTP/JSON* API as an interface for setting or retrieving values from the store. Other components watch the state of the *etcd* store to bring themselves up to the desired state. Data being stored in the *etcd* store are deployed services, pods and replication information.

#### 2.7.3. Scheduler

The scheduler component is responsible for the deployment of pods and services on the cluster nodes. The scheduler has the information about the availability of resources on a node and schedules unscheduled pods on the nodes accordingly. Along with scheduling, the scheduler also tracks resource utilization of each node and ensures that workload scheduled is not in excess to the resources available [15].

#### 2.7.4. Controller-manager

The controller manager is the process embedding the different types of controllers like the Replication Controller or the DaemonSet Controller on a kubernetes master. The controllers query the API server to manipulate the resources like pods and services which they manage.

## 2.8. Worker Node

The worker node also called as minion node is where the containers are actually deployed. The worker contains all the necessary services needed to manage the networking between containers, communicate with the master node and assign resources to the scheduled containers [16]. Every worker node must run the container runtime i.e docker and other components stated below to ensure proper communication with the master.

#### 2.8.1. Docker

Docker runs on each of the worker nodes. It is responsible for downloading the docker images and running the configured pods by starting the container.

#### 2.8.2. Kubelet

*Kubelet* gets the pod definition from the api-server and is responsible for maintaining the pod in the desired state. *Kubelet* is the worker service that monitors the health of each pod and communicates the status of each node via a heartbeat message to the master. If the pod is not in the desired state, it is redeployed to the same node [15]. *Kubelet* is also responsible for communicating with the *etcd* storage to get information about the services and update the storage about newly created ones.

#### 2.8.3. Kube-Proxy

*Kube-proxy* acts as a network proxy and a load balancer. It is responsible for networking of *TCP* and *UDP* packets to the appropriate container based on the IP address of each packet [15] [16].

#### 2.8.4. Kubectl

*kubectl* is a command line tool that communicated with the API server to fetch important information about the nodes, pods, services and events in the cluster.

## 3. DESIGN

### 3.1. About the Application

We are using a series of machine learning algorithms and shall compare their performance in terms of classifier accuracy. To build our classifier and check its accuracy, we shall partition the data into 2 subsets - training data and classification data. We shall choose one of the subsets for training and other for testing and predicting classifier accuracy.

### 3.2. The dataset

We are using the SMS Spam Collection dataset from UCI's public machine learning repository to build our classifier. The dataset consists of 5574 text messages, with each classified as a spam or ham message [19].

### 3.3. Algorithms Used

#### 3.3.1. SVM

SVM (Support Vector Machines) are discriminative classifier models that classify the data as belonging to one of the two possible classes, in our case class ham and class spam. The SVM

algorithm attempts to categorize data by drawing a hyperplane between the data points such that points on either end of the hyperplane belong to one of the two possible classes. An SVM model is representation of given data points in space such that the division between the classes of data is as clear as possible [20]. The aim of the algorithm is to find the optimum hyperplane that defines the largest possible margin between the classes of data.




Fig. 3. The maximum margin hyperplane for SVM [21]

### 3.3.2. SVM Kernels

A simple SVM transformation works for data with linear decision boundaries, where the different data points lie on either side of the decision boundary. But for scenarios where a linear boundary is efficient to classify data, we need to transform the data from its original coordinate space into new space so that a linear decision boundary can be drawn to separate the data in this transformed space [22]. The kernel is a similarity function. It is a method to compute similarity between attributes in the transformed attribute space. We shall use svm kernel methods as they eliminate the curse of dimensionality problem as they perform the computations in the original attribute space [22]. We shall avoid mathematical details about SVM and the kernel as they are beyond the scope of this report.

**Polynomial kernel** - Polynomial kernel, as the name suggests, looks not only at the given features of the input samples but at their combination [23]. If a pair of words give interesting information rather than the individual words alone, we can use a quadratic kernel. Similarly, for occurrences of triplets of words, we can use a cubic kernel. Linear kernels are special case of polynomial kernels where the quadratic factor is 1.

**Radial kernel** - Radial basis kernel maps the data into infinite dimensional space to extract relations between the variables. It helps us to draw a circular decision boundaries to pick dependent features.

### 3.3.3. K - Nearest Neighbor

The k-nearest neighbor algorithm classifies a data point as belonging to either output class by taking a majority vote amongst its k-nearest neighbors. The algorithm computes distances between the test data point and all the stored data points or neighbors. *Euclidean distance* [24] is used to compute distance metrics between attribute values of the neighbors. A sample figure explaining the algorithm is given below.

There are 2 classes of data items in the figure - circles and squares. We have an unknown data point - the star and wish to find its real class using k-nearest neighbor algorithm. For  $k=3$ , the




Fig. 4. An example for the K-Nearest Neighbor algorithm [25]

3-nearest neighbours according to the distance metrics are all circles. Hence, the unknown data point is then classified as a circle. This assumption changes for higher values of  $k$ . To eliminate the effect of noise and outliers on the classification result, a weight is associated with each data point, with closer neighbors having higher weight over the distant neighboring points.

### 3.3.4. AdaBoost

AdaBoost (Adaptive Boost) works on the principle of generating better classifiers by adapting existing classifiers according to the data over which the classification process is to be performed. AdaBoost combines the output of several classifiers into a weighted sum that represents the final output of the algorithm [26]. At each step in the algorithm, a boosting sequence is performed that finds weak classifiers by comparing the ratios of misclassified samples by each classifier and then tweaks them to generate classifiers that provide more accurate data. Unlike SVM, AdaBoost works with only those attributes of the data that improve the predictive power of the classifier and does not suffer from the *curse of dimensionality* [27] drawback of most data mining algorithms.

### 3.3.5. Naive Bayes Classifier

Naive Bayes classifier works on the principle of conditional probability. Bayes classifier assigns each sample a probability of belonging to one class or another. The classifier shall maintain a bag of words along with the count of each word occurring in the spam messages. This word count shall be used to calculate and store the word probability in a table that shall be cross-referenced to determine the class of the record on classification data [28].

A selected few words have more probability of occurring in a spam messages than in the legitimate ones. Eg: The word "Lottery" shall be encountered more often in a spam message. The classifier shall correlate the bag of words with spam and non-spam messages and then use Bayes Theorem to calculate a probability score that shall indicate whether a message is a spam or not. The results shall be verified with the results available on the training dataset and the classifier accuracy shall be calculated. The classifier shall use the Bayesian theorem over the training dataset to calculate probabilities of such words that occur more often in spam messages and later use a summation of scores of the occurrence of these word probabilities to estimate whether a message shall be classified as spam or not. After working on several samples of the training dataset, the classifier shall have learned a high probability for spam based words whereas, words in legitimate message like family member or friends names shall have a very low probability of occurrence.

Once the training process has been completed, the posterior probability for all the words in the new input email is computed using Bayes theorem. A threshold value shall be defined to classify a message into either class. A message's spam probability is computed over all words in its body and if the sum total of the probabilities exceeds the predefined threshold, the filter shall mark the message as a spam [29].

### 3.4. The Classification process

#### 3.4.1. Building the classification model

To address the problem of incoming spam messages, a model shall be developed using the various classification techniques to correctly classify each incoming email/text message as a spam or a legitimate one. The training dataset to build the model consists of 5574 message records.

To develop an efficient training model, we shall partition the data into 2 subsets - the training data and testing data. We shall choose one of the subsets for building a classification model and the other to evaluate the model's performance. The training process involves creating a document term matrix for the occurrence of each word in the spam marked messages. This matrix describes the frequency of terms in each of the messages, with the rows corresponding to documents in the dataset and columns corresponding to the terms. Using the document term matrix, we find highly repeated words in spam messages across the training set. We vectorize the training and the test data set using these word frequencies and provide this vectorized data as input to the various algorithms in our program.

#### 3.4.2. Testing the model

Each algorithm used in the application uses its own strategy to build a classification model from the training set examples. The model's performance is evaluated by using it to classify the test dataset. Using the classification results, a confusion matrix is built to analyze the algorithm's performance on the test dataset. The confusion matrix specifies what number of spam messages were correctly classified as spam or incorrectly classified as ham and vice-versa. Based on these results, we calculate the accuracy of classification by dividing the number correct results by the total records in the dataset.

Further, a higher classification accuracy shall be achieved through filtering by looking at the message header i.e the sender's number/name. Thereby, if a message from a particular sender is repeatedly marked as spam by the user, the classifier need not evaluate the message body if it is from the same sender.

## 4. DEPLOYMENT

Our application will be deployed using Ansible playbook. Automated deployment is performed on a multi-node cluster on both Chameleon and Kilo clouds. Deployment script shall install all required software components along with the project code to Kubernetes cluster nodes using the Docker image.

### 4.1. Dockerizing the application

To containerize our application, we need to create a docker image for it. To create a docker image we need to create a docker file that describes the semantics to create the docker image. Docker version used for this process is 1.12.6.

#### 4.1.1. The Dockerfile

Docker can build images automatically by reading instructions from a Dockerfile. A dockerfile is a series of text commands that

a user can run on a command line to build an image [30]. Instructions within a Dockerfile have the following format:

INSTRUCTION arguments

**Fig. 5.** Syntax of every command in a dockerfile

Docker reads the instructions from a dockerfile in the order they are written. The very first line of a docker file must be a 'FROM' clause which specifies the base image from which the image is being built. As our application works with R, we need to specify the base package as one that shall help us run our R application. *DockerHub* [31] has a r-base package that binds the latest version of R and its libraries together [32]. The argument following the FROM clause is a repository/tag name that the docker engine automatically looks up on dockerhub.

```
FROM r-base
COPY . ~/myscripts
WORKDIR ~/myscripts
RUN Rscript -e "install.packages('tm')"
RUN Rscript -e "install.packages('e1071')"
RUN Rscript -e "install.packages('RWeka')"
RUN Rscript -e "install.packages('ada')"
RUN Rscript -e "install.packages('rbenchmark')"
CMD ["Rscript", "spamdetection.r"]
```

**Fig. 6.** Dockerfile of the application

Following the FROM statement, we can specify all actions we need to perform like specifying the work directory for building the image or copying files into our work directory so that every resource is available under a single folder. Our application needs several text mining packages to be available as a part of the installation and thus we use the docker RUN command to download these dependencies and make them part of the docker image. The R library in the docker container shall pre-install all the dependencies as the image is built and we shall only care about running the application logic. Lastly, our docker file ends with a CMD statement that specifies the command that the container is supposed to execute on successful instantiation. CMD accepts an array of command names followed by the parameters. Since we want our image to execute our R application we shall specify the command as stated in the dockerfile below.

#### 4.1.2. Creating a docker image

Once the Dockerfile is built, we can build the image using docker build command as follows.

The built image is placed in our machine's local docker registry and can be viewed with the 'docker images' command.

REPOSITORY	TAG	IMAGE ID	SIZE
spamdetectionapplication	latest	a4814eeadf	646.1 MB
r-base	latest	16fe32463d	645.6 MB

**Fig. 8.** Listing images in the local repository

#### 4.1.3. Running the application

Once our docker image is built, it can be executed using the docker run command.

```
cc@rahpsing-055:~$ sudo docker build
-t spamdetectionapplication ~/dockerDirectory/

Sending build context to Docker daemon 12.8 kB
Step 1 : FROM r-base
latest: Pulling from library/r-base
c560cd7bd403: Pull complete
b190a5321f19: Pull complete

Digest: sha256:e192edf861d61caff0b329436...
Status: Downloaded newer image for r-base:latest
---> 16fe32463daa

Step 2 : COPY . ~/dockerDirectory/
---> a6717da2ec47
Removing intermediate container 8fb4aaaec0d8

Step 3 : WORKDIR ~/dockerDirectory/
---> Running in 0220b193ed11
Removing intermediate container 0220b193ed11
Step 4 : RUN Rscript -e "install.packages('tm')
      <installation-logs>
      ..... <intermediate-steps>.....

Step 9 : CMD Rscript spamdetection.r
---> Running in 4edb44f2ffb6
Removing intermediate container 4edb44f2ffb6
Successfully built a4814eeadf35
```

**Fig. 7.** Building an image line by line from a dockerfile

```
docker run spamdetectionapplication
```

**Fig. 9.** Running a docker image

We can check the container id of our application along with other important information using the 'docker ps' command.

#### 4.1.4. Sharing the docker image

By default, the docker CLI points to docker's public registry which is located on dockerhub's website. We need to create a docker account to upload our application image so that it could be directly referenced by kubernetes later. Once a docker account is registered, create a public repository and give it a name. The repository shall be accessible by the format '<username>/repositoryname'.

To tag the current docker image that we created on our system we need to inform the docker engine to point to our registry. To do so, we shall use login command provided by the docker CLI.

```
docker login
```

**Fig. 10.** Linking a dockerhub account to the local registry

The login request shall ask for a valid docker account userid and password. Once authenticated, the docker engine shall establish a link with the registry associated with the username. To upload our docker image to the repository, we need to tag the image. Docker provides the 'tag' command to do so.

```
docker tag image_name <username>/<repository_name>:
<tag>
```

**Fig. 11.** Tagging a local docker image

The <tag> can be any name given by the user to uniquely identify the image. Once tagged, we need to push the image to the repository using the command line 'docker push' command.

```
docker push rahpsing/kubernetesi524:spamdetection
```

**Fig. 12.** Pushing an docker image to a registry

As the image is now publicly available on the repository, we can do a simple 'docker pull' from any machine to deploy the container with no dependencies. This is possible as the image wraps together the application code file, the data file and dependent r-base package from the latest version of R available on dockerhub's library.

## 4.2. Deployment via kubeadm

Kubeadm is a part of kubernetes 1.4 distribution that allows users to install and set up a Kubernetes cluster. Kubeadm works with local VM's, cloud servers or physical servers [33].

### 4.2.1. Creating a cluster

To create a kubernetes cluster, it is essential to install docker, kubelet, kubectl and kubeadm on all the machines that are to be a part of the network. Kubelet shall help to start pods and containers on all the machines of the network. Kubectl allows us to monitor the activities of the cluster once it is up and running. Primarily, it is used only on the master node. Kubeadm is used to setup the cluster by allowing multiple worker nodes to bind with the master on a unique network identifier token.

Once the mentioned components are installed, we need to initialize the master so that it can accept requests from other service nodes to join the cluster. Master initialization is done with the 'init' command.

```
kubeadm init
```

**Fig. 13.** The master initialization command

Executing the above command tells kubernetes that the host machine shall serve as the master in the cluster. Kubeadm initializes all the other dependent components like the API server and generates a unique key that identifies the master in the network. Client nodes can issue a join request with the master by using the unique identifier as a part of the join command.

```
kubeadm join <unique_token>
```

**Fig. 14.** Syntax of the join command

A few seconds after running the join command we can query the master to list the available nodes in the cluster using 'kubectl get nodes' command.

```
kubeadm join --token 87ce11.50ab6a5eea 192.168.0.204
$ kubectl get nodes
NAME                STATUS    AGE
rahpsing-056        Ready     2m
rahpsing-057        Ready     2m
```

**Fig. 15.** Registering nodes to the cluster

Before the master is ready to schedule pods, it is imperative that the API server and the kubedns service are up and running. We can check the status of all the services of the system using the 'get nodes' command with the '-all-namespaces' argument.

The kubedns service is responsible for networking with its client nodes. Without the kubedns service, the kubernetes master shall be unable to schedule pods on the worker nodes as pods cannot communicate with each other. To get the service up and running, we need to install a pod network. Kubernetes provides many add-on services that can be used to setup the cluster network policy and enable networking. Of the list of available add-ons we shall use 'weave-net' as it provides us with a network policy and does not require an external database [34].

```
kubect1 apply -f https://git.io/weave-kube
```

**Fig. 16.** Creating the weave-net pod-network

```
$ kubectl get pods --all-namespaces
NAME                                READY    STATUS
etcd-rahpsing-056                   1/1      Running
kube-apiserver-rahpsing-056         1/1      Running
kube-controller-manager-rahpsing-056 1/1      Running
kube-discovery-2849056221-5xws4     1/1      Running
kube-dns-2247936740-m9198           3/3      Running
kube-proxy-amd64-d91cy              1/1      Running
kube-proxy-amd64-ycg5e              1/1      Running
kube-scheduler-rahpsing-056         1/1      Running
weave-net-1sk03                     2/2      Running
weave-net-p9cpd                     2/2      Running
```

**Fig. 17.** Services running in the cluster

#### 4.2.2. Creating a pod

A pod as described initially is the smallest unit of work in a kubernetes cluster. Semantics of a pod are defined in a pod.yaml file which is then passed to the kubectl CLI to initialize the pod and maintain its desired state.

The pod.yaml file specifies the containers that compose a pod. Each container works on top of a docker image that it runs once instantiated. We shall use the spamdetection image for our container that we created by dockerizing our application and pushing the image to the repository. By default, kubernetes searches for image tags specified in the pod.yaml file on docker's public repository. If the matching image is found, the pod is created successfully or it results in a pod failure.

Once the semantics of a pod are fixed we can use kubectl to instantiate the pod. We can also use 'kubectl apply' if we wish to make changes to the definition of the pod and want the implementation to reflect it.

kubectl provides 'get XXX' commands to query status of any type of object in the system. We can check the status of our pods

```
kubect1 create -f pod.yaml
```

**Fig. 18.** Creating a pod

```
$ kubectl get pods
NAME                READY    STATUS    AGE
spamdetectionimage  0/1      Running   46s
```

**Fig. 19.** Status of pods in the cluster

by using the API's get pods command.

To get a detailed view of each element we can use the 'describe <element\_name>' command where <element\_name> is unique under each type of component in the system.

#### 4.2.3. Creating a deployment

Similar to creating a pod, a deployment can be created in Kubernetes using the same yaml file syntax. The only notable difference in this case is the value of the 'type' and the 'apiVersion' field in the yaml file. The 'type' field is set to 'deployment' in case of creating a deployment while the apiVersion field is set to 'extension/v1beta1' as version 'v1' which was used for pod creation does not support deployments. With the deployment file, we have the flexibility of specifying the number of replicas we want to deploy. We set the number of replicas to 2, which means that kubernetes will ensure that at any given point in time the system will have 2 replicas of our application running. After defining the number, we define the containers whose replicas we wish to maintain.

kubect1 command to create a deployment is the same as the one we used to create a pod. The Kube Control API parses the yaml file and understands the type of object it is expected to create.

As with pods, the command to list the deployments running on the cluster is 'kubectl get deployments'. Now that our deployment is created we can check its status by using various kubectl commands. As we set the number of replicas to be 2, there should be atleast 2 pods running our application which can be verified as below.

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: spamdetectionapplication
spec:
  replicas: 2
  template:
    metadata:
      labels:
        app: echo
    spec:
      containers:
        - name: container2
          image: rahpsing/kubernetesi524:spamdetection
            application
          ports:
            - containerPort: 80
```

**Fig. 20.** Deployment file of the application



```
$ kubectl get pods
NAME                                READY    STATUS
spamdetectionapplication-1         1/1      Running
spamdetectionapplication-2         1/1      Running
```

**Fig. 21.** Deployment with 2 replicas

#### 4.2.4. Fetching the application output

As stated previously, a pod is a collection of multiple related containers. Our application executes in of the containers within the pod. Since our application prints output to the console, we shall be able to view it by accessing logs of the relevant pod.

```
kubectl get pods | grep spamdetectionapplication*

kubectl logs <pod_name>
```

**Fig. 22.** Fetching logs for the application

As each pod is hosted on a separate node with a unique ip address, we can issue a *curl* or a *wget* request to the node to fetch the output. Since the output produced by our application spans over multiple lines we shall redirect and save the output to a file so that it can be referenced at any point of time with ease.

```
curl https://<node_ip>:<node_port>/containerLogs/
default/<pod_name>/spamdetectionapplication
--insecure > output.txt
```

**Fig. 23.** Saving container logs to a file

Following the above steps, the application output shall be available in the file 'output.txt' in the current directory.

#### 4.2.5. Scheduling pods on the master

By default, kubernetes does not allow pods to be scheduled on the master node for security purposes. However, this setting can be overridden by removing the taint on the master node and allowing pod execution.

```
kubectl taint nodes --all node-role.kubernetes.io/
master-
```

**Fig. 24.** Command to schedule pods on the master

Execution of the above statement shall remove the taint 'node-role.kubernetes.io/master' from all the nodes in the network including the master node and thus allowing the scheduler to schedule pods across the network [33]. This feature also allows us to create a single node cluster.

## 5. DEPLOYMENT USING CMD5 MODULE

All the steps required for kubernetes have been implemented using Ansible along with python's CMD5 module [35]. Python's CMD module is a very useful package for creating shell like commands but it does not offer the extension of dynamic integration with any new customized commands which was the reason why CMD5 module was introduced. CMD5 has been rewritten on earlier versions of CMD3 and cloudmesh client.

### 5.1. Installation of CMD5 Module

The most convenient way to install CMD5 module is to checkout all its necessary repositories and install the setup file so that we may create new commands for our use.

Use the terminal on your windows or mac workstation and clone the following 3 repositories. It is mandatory to have *Git* [36] already installed.

```
$ git clone
https://github.com/cloudmesh/cloudmesh.common.git
```

**Fig. 25.** Cloning the cloudmesh.common repo

```
$ git clone
https://github.com/cloudmesh/cloudmesh.cmd5.git
```

**Fig. 26.** Cloning the cloudmesh.cmd5 repo

```
$ git clone
https://github.com/cloudmesh/cloudmesh.sys.git
```

**Fig. 27.** Cloning the cloudmesh.sys repo

After cloning these repositories, we need to enter each repository and execute the following statements.

```
$ python setup.py install
```

**Fig. 28.** Executing the setup file

```
$ pip install .
```

**Fig. 29.** Installing the command

These commands will install common, sys and cmd5 modules in your workstation which would lead to successful installation of the CMD5 module.

### 5.2. Running the CMD5 module

To launch CMD5, you should open the terminal and type in cms. This would launch the CMD5 module with few already existing commands like EOF, echo, help, stopwatch etc.

### 5.3. Creation of customised CMD5 commands

To create your own customized CMD5 command, you must type in the following.

```
$ cms sys command generate <command_name>
```

**Fig. 30.** Creation of customised CMD5 command

This command would create a CMD5 command with the name provided as an argument. It downloads a template from



cloudmesh called cloudbmesh.bar and generates a new directory cloudmesh.<command\_name> with all the files which are needed to create the new command. It also registers it dynamically with cloudmesh. To finish we need to move into the newly created directory and execute the setup and install commands.

## 6. CLOUDMESH.KUBERNETES COMMAND

For our kubernetes deployment, we have created our own customized kubernetes command called cloudmesh.kubernetes. We have extended the kubernetes command to create various sub-commands that help us in setting up a cloud cluster, deploying kubernetes and performing benchmarks of the application.

We have created 8 Kubernetes CMD5 sub-commands and each of those commands trigger execution of relevant ansible scripts which perform the automated deployment for our spam detection application.

### 6.1. Kubernetes Name command

```
$ kubernetes name <cluster_name>
```

**Fig. 31.** Command to give the cluster a name

The name option assigns a name to the cluster that is passed as an argument. The cluster when deployed will be identified by this name on the cloud.

### 6.2. Kubernetes Size command

```
$ kubernetes size <size_of_cluster>
```

**Fig. 32.** Command to define the cluster size

The size option creates a cluster with the desired number of nodes. For example, Kubernetes size 2 would create a cluster with 2 nodes on the specified cloud.

### 6.3. Kubernetes Image command

```
$ kubernetes image <image_name>
```

**Fig. 33.** Command to define the type of image for the cluster

This command is used to define the type of image to be used by the cluster. For example, the image name could be CC-Ubuntu16.04-20160610 which is a valid chameleon image. We can also define any valid image for the kilo cloud as well. The hard restriction here is that whichever cloud we decide to use for the cluster, the image name specified to be used should be a valid image available on the cloud.

### 6.4. Kubernetes Flavor command

```
$ kubernetes flavor <flavor_size>
```

**Fig. 34.** Command to define the flavor for the cluster

This command is used to define the flavor for each instance within the cluster. The 3 different types of flavors can be m1.small, m1.medium and m1.large. Depending upon the small, medium or large flavor value we are allocated an instance with the appropriate configurations.

### 6.5. Kubernetes Cloud command

```
$ kubernetes cloud <cloud_name>
```

**Fig. 35.** Command to define the cloud on which the cluster will be created

This command is used to specify the cloud on which the cluster is to be created. Deployment of our application was performed using kilo and chameleon clouds as parameter to this command. This setting will create the cluster on the desired cloud when the kubernetes cluster deploy command will be executed.

### 6.6. Kubernetes cluster info command

```
$ kubernetes cluster info
```

**Fig. 36.** Command to get details of the cluster to be created

This command would list all the details of the kubernetes cluster which was set using the name, size, image, flavor and cloud commands. This command helps us to know all the details for our kubernetes cluster. This command would print out the name of the cluster, its size, the image, the flavor used and the cloud on which deployment is expected to be performed.

### 6.7. Kubernetes cluster deploy command

```
$ kubernetes cluster deploy
```

**Fig. 37.** Command to deploy the cluster on the cloud

This command would deploy the Kubernetes cluster with the details set in the name, size, image, flavor and cloud commands. If any of the above parameters are not set, this would not deploy the cluster and ask the user to enter the required details before deploying it.

The deploy command first creates a cluster with the number of instances equal to the size specified using the size command with the image and flavor specified and deploys the cluster on the cloud specified using the kubernetes cloud command. Once the cluster is created, the command invokes a series of ansible tasks which install the deployment specific pre requisite software on these instances, make one instance as the master and others as slaves. Next, Ansible tasks are invoked which install kubernetes on all the instances and joins the slaves to the master creating a kubernetes cluster.

### 6.8. Kubernetes cluster benchmark command

```
$ kubernetes cluster benchmark
```

**Fig. 38.** Command to benchmark the kubernetes cluster

This command would benchmark the kubernetes cluster for our spam detection application. This command downloads our application docker image hosted on dockerhub and deploys it to the cluster. Kubernetes creates the desired number of pods based on our deployment specification and executes our application by following commands in the dockerfile. The kubernetes master is responsible for scheduling of pods for the deployment. The pods can begin execution only when all the network services

in the cluster are ready. We have used *weave-net* as our pod-network provider. Weave-net takes a few minutes to initialize its networking services and hence we have to wait for the pods to be ready before we can fetch the application output.

## 7. ANSIBLE DEPLOYMENT

Ansible is used for automated deployment of our application on any local or remote cluster.

### 7.1. Inventory file

The inventory file is used to help ansible know which all instances it needs to connect to along with a list of the required tasks. For our kubernetes deployment, the IP addresses which are generated by the respective cloud are stored in the inventory file using *ssh-keyscan* command. For our application, we create an inventory file having one instance allocated to the kubernetes-master host and all other instances allocated to the kubernetes-slave hosts. The number of instances depend upon the kubernetes size command.

### 7.2. Roles and Playbook

The roles concept in Ansible allows us to draw a logical meaning to the tasks being performed. We can create as many roles as possible and define the tasks in those roles. We can then define the list of the role-names in the yaml file which would be executed.

For our project, we have created 5 yaml files which serve as our playbooks and invoke ansible roles that perform tasks on the remote cloud cluster. Each yaml file contains the host name and the roles it shall invoke on execution.

The first playbook that is triggered when the CMD5 kubernetes cluster deploy command is executed is the *installations.yml*. This playbook performs initial cleaning of previously present files if any on the cluster, installs git and performs an update operation on all the instances in the cluster.

The second playbook *kubernetes.yml* invokes the setup role which installs kubernetes and its dependencies on all the nodes in the cluster.

The third playbook to be executed is *master.yml* playbook. It initializes one of the nodes in the cluster as master and initializes the kubernetes cluster.

After the master is up and running, the *slaves.yml* playbook is executed which connects the other instances in the cluster to the master instance. Each of the slaves executes the *kubeadm join* command which consists of the master's ip address and a unique connection token generated by the master to identify itself in the cluster. The first 4 playbooks are part of the Kubernetes deploy command.

Once the kubernetes cluster is setup and the master and slave nodes are initialized we execute the *runningapplicationonkubernetes.yml* playbook which deploys the spam detection application on one of the slave which is dynamically chosen by the master and performs benchmarking of the time taken by various algorithms in classifying an email as spam or valid.

## 8. CLASSIFICATION RESULTS

After running all the aforementioned algorithms on our data file the following results were achieved.

Algorithm	Classification Accuracy (%)
SVM Linear Kernel	93.58
SVM Polynomial Kernel	92.12
SVM Radial Kernel	93.72
Naive Bayes Classifier	91.76
Ada Boost	93.9

Fig. 39. Classification results

## 9. ISSUES TACKLED WITH KUBERNETES SETUP AND DEPLOYMENT

### 9.1. ImagePullBackOff

```
$ kubectl get pods
NAME          READY   STATUS             RESTARTS
myapplication  0/1     ImagePullBackOff   21
```

Fig. 40. The ImagePullBackOff error

The ImagePullBackOff error indicates that the container was unable to fetch the desired image from the repository specified. This happens in case of improper tag names associated with the application image or reference to a repository that doesn't exist. By default, any text specified against the 'image' key in the pod or deployment file makes a direct lookup on dockerhub. Thus any docker image that has been built on the system cannot be referenced by the kubernetes engine unless it is tagged and uploaded to a public repository on dockerhub. Dockerhub allows users to create public repositories for no charges. Thus, the only way to use your custom created image is to register for an account on dockerhub, create a public repository, upload your local image with a unique tag name and use the repository and tag name in the yaml file. This issue took a lot of our development time as there is no clear documentation which explicitly describes the kubernetes interpretation of the value of the 'image' key specified in the pod yaml file.

### 9.2. CrashLoopBackOff

```
$ kubectl get pods
NAME          READY   STATUS             RESTARTS
myapplication  0/1     CrashLoopBackOff   49
```

Fig. 41. The CrashLoopBackOff error

Kubernetes allows us to execute instructions after a container has been deployed successfully. The output of these instructions is visible in the container's logs which can be accessed by using the 'get logs' primitive of the kubectl CLI. Commands to be executed can be specified under the 'args' key of the container

specification in the pod definition file. However, kubernetes defines strict parsing for those arguments with no pointers to indicate where the execution failed. Any issues with the yaml file shall trigger a crashloopbackoff error with no traces of line numbers or clear log messages to guess what went wrong. It was after considerable research and trial and error that we figured out that an extra space between arguments to the sleep command was causing the error.

### 9.3. Kube DNS service not working

```
$ kubectl get pods --all-namespaces
NAME                                READY  STATUS
etcd-rahpsing-056                  1/1    Running
kube-apiserver-rahpsing-056        1/1    Running
kube-controller-manager-rahpsing-056 1/1    Running
kube-discovery-2849056221-5xws4    1/1    Running
kube-dns-2247936740-m9198         0/3    Container
                                      Creating
```

**Fig. 42.** Failure of the kube-dns service

Kubedns is an integral part of the kubernetes master setup as it enables pods to communicate with each other. Before creating any pods or deployments we need to setup a pod-network. The pod network defines security policies and establishes networking within the cluster. Kubernetes provides various addons to be used as its pod network. We tried to use *flannel* but were unable to bind it to our network and get the dns service working. After numerous attempts we figured out that *weave-net* is the best pod-network that we could use as its network was updated to work with the *CNI* networking policy introduced in Kubernetes 1.4. Also, using weave net as our pod network allows us to use the weave-kube service where we can observe our cluster components and their configurations on a GUI accessible by a web browser.

### 9.4. Unsuccessful download requests from within a kubernetes pod

Initially, our R application program started with many download and install-package statements as we needed several text mining packages to implement the respective algorithms. Even after successful deployment, we couldn't find the application output in the container logs. It took us a fair amount of time to figure out that something wasn't working correctly as the pod logs failed to highlight the issue. To debug in deeply, we started deploying applications with simple print statements to check their behavior. After several attempts, we could draw a conclusion that deployment would result in a failure if the program included package import statements that needed to be downloaded from an external repository. We ran the same application image with docker only to find it working perfectly. Thus, we could state that, from within a kubernetes pod, an application is unable to reach the network and download the desired packages. We could find no precise explanation of this behavior on the kubernetes forum or issues opened on stack overflow or kubernetes git page.

As a workaround, we segregated the download statements from the R program and included it within our docker file. This modification would download and install the R libraries and packages as a part of the build process of the docker image, eliminating the need to install it post deployment. With the dependent libraries as a part of the docker container, the application program


would run successfully and generate the output as a part of the logs for the pod it is associated with.

## 10. BENCHMARKING RESULTS

Application benchmarking was accomplished by executing the application on a kubernetes cluster with two nodes on both chameleon and kilo clouds. To observe execution time of different algorithms, we inserted checkpoints in the R application to measure time taken by each algorithm on different clouds. We referred to the application output file to get the runtime of all 5 algorithms on our dataset.

Algorithm	Chameleon cloud execution time (seconds)	Kilo cloud execution time (seconds)
SVM Linear Kernel	1.15	1.22
SVM Polynomial Kernel	1.68	1.72
SVM Radial Kernel	1.75	1.78
Ada Boost	20.25	21.75
Naive Bayes	5.93	6.21

**Fig. 43.** Algorithm execution time on Chameleon and Kilo cloud

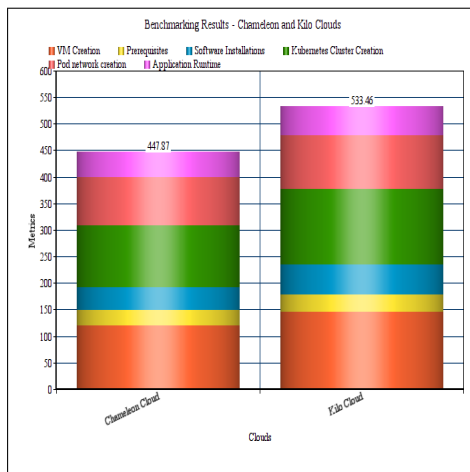


**Fig. 44.** Graph plot of algorithm execution time on chameleon and kilo clouds

Further, the stopwatch command on CMD5 module was used to benchmark execution time of kubernetes cluster creation, software deployment and installation on both chameleon and kilo cloud. By segmenting the all-round application runtime into software installation, cluster creation and code deployment we can clearly observe the performance of the cluster on each cloud. The benchmarking time mentioned above is an average of 5 runs on both chameleon and kilo clouds. Segregating the application deployment process into several tasks allowed us to compare time taken by each step on both clouds.

Tasks	Chameleon cloud execution time (seconds)	Kilo cloud execution time (seconds)
VM creation	122	148
Prerequisites	28.23	38.21
Software installation	42.83	56.43
Kubernetes Cluster Creation	20.25	21.75
Pod network installation	5.93	6.21
Application runtime	47.28	53.24

**Fig. 45.** Algorithm execution time on Chameleon and Kilo cloud



**Fig. 46.** Graph plot of application deployment time on chameleon and kilo clouds

**VM creation** - This task included cloudmesh and kilo cluster creation.

**Prerequisites** - Installation of git and updating the system using apt-get update.

**Software installation** - This step included downloading docker, kubeadm and kubectrl packages, extracting and installing them.

**Kubernetes Cluster Creation** - This step included creating a kubernetes cluster using the kubectrl init command. The init command takes some time to generate a unique token for the master machine which can be used by clients to join the cluster and serve as slaves. This step accounts for the time taken for clients to join the cluster.

**Pod network installation** - Pods and nodes in a kubernetes cluster cannot communicate without agreeing on a common pod network. We used weave-net as the pod network for our cluster. Pods cannot start running services unless all the services that are part of the CNI are active.

**Application runtime** - This is the time taken by the entire pro-

gram to run including reading of the data file and execution of each algorithm.

## 11. SUMMARY

We have successfully deployed and benchmarked a standalone application on a kubernetes cluster on both Chameleon and Kilo Clouds. This project helped us understand the notion of application portability using docker and deployment management using kubernetes. Benchmarking results proved that Chameleon cloud performed far better than Kilo cloud in all aspects of application deployment as well as execution results.

While the application code was developed in a reasonable time, a lot of effort was put to find the correct steps of a standalone application deployment with kubernetes. We have highlighted a few major issues that we faced in this report. Because we went through a barrage of failures, we ended up discovering more about kubernetes and its deployment features than what was needed for this project. However, our struggles also convey that although kubernetes offers a variety of cluster management functionalities, it has a steep learning curve. Any organization that intends to switch to kubernetes for deployment management would need to train its developers and use it on sandbox environment for a while before implementing it in production.

## 12. ACKNOWLEDGEMENT

We would like acknowledge our professor Gregor von Laszewski and all the associate instructors for motivating and guiding us throughout this project.

## REFERENCES

- [1] "R:what is R?" Web Page, accessed: 2017-04-02. [Online]. Available: <https://www.r-project.org/about.html>
- [2] "What is docker?" Web Page, accessed: 2017-04-02. [Online]. Available: <https://www.docker.com/what-docker>
- [3] "Kubernetes - production-grade container orchestration," Web Page, accessed: 2017-04-11. [Online]. Available: <https://kubernetes.io/>
- [4] "Ansible, deploy apps. manage systems. crush complexity," Web Page, accessed: 2017-04-12. [Online]. Available: <https://www.ansible.com/it-automation>
- [5] "ssh protocol," Web Page, accessed: 2017-04-12. [Online]. Available: <https://www.ssh.com/ssh/protocol/>
- [6] Microsoft, "Windows remote management (windows)," Web Page, accessed: 2017-10-22. [Online]. Available: [https://msdn.microsoft.com/en-us/library/aa384426\(v=vs.85\).aspx](https://msdn.microsoft.com/en-us/library/aa384426(v=vs.85).aspx)
- [7] "Ansible architecture," Web Page, accessed: 2017-04-12. [Online]. Available: [http://docs.ansible.com/ansible/latest/dev\\_guide/overview\\_architecture.html](http://docs.ansible.com/ansible/latest/dev_guide/overview_architecture.html)
- [8] Wikipedia, "Ansible(software) - wikipedia," Web Page, accessed: 2017-05-19. [Online]. Available: [https://en.wikipedia.org/wiki/Ansible\\_\(software\)](https://en.wikipedia.org/wiki/Ansible_(software))
- [9] "Inventory - ansible documentation," Web Page, accessed: 2017-10-19. [Online]. Available: [http://docs.ansible.com/ansible/latest/intro\\_inventory.html](http://docs.ansible.com/ansible/latest/intro_inventory.html)
- [10] "Playbook," Web Page, accessed: 2017-04-12. [Online]. Available: <http://docs.ansible.com/ansible/playbooks.html>
- [11] Wikipedia, "Virtualization - wikipedia," Web Page, accessed: 2017-04-12. [Online]. Available: <https://en.wikipedia.org/wiki/Virtualization>
- [12] M. Rouse, "What is application containerization (app containerization)?" Web Page, accessed: 2017-04-18. [Online]. Available: <http://searchitoperations.techtarget.com/definition/application-containerization-app-containerization>
- [13] Wikipedia, "Virtual machine - wikipedia," Web Page, accessed: 2017-04-12. [Online]. Available: [https://en.wikipedia.org/wiki/Virtual\\_machine](https://en.wikipedia.org/wiki/Virtual_machine)

- [14] O.S.Tezer, "Docker explained: Using dockerfiles to automate building of images," Web Page, Dec 2013, accessed: 2017-04-24. [Online]. Available: <https://www.digitalocean.com/community/tutorials/docker-explained-using-dockerfiles-to-automate-building-of-images>
- [15] Wikipedia, "Kubernetes-wikipedia," Web Page, accessed: 2017-04-18. [Online]. Available: <https://en.wikipedia.org/wiki/Kubernetes>
- [16] N. Isabekyan, "Introduction to kubernetes architecture," Web Page, Jul 2016, accessed: 2017-04-24. [Online]. Available: <https://x-team.com/blog/introduction-kubernetes-architecture/>
- [17] J. Ellingwood, "An introduction to Kubernetes," Web Page, accessed: 2017-04-18. [Online]. Available: <https://www.digitalocean.com/community/tutorials/an-introduction-to-kubernetes>
- [18] K. Marhubi, "Kubernetes from the ground up:the API server," Web Page, accessed: 2017-04-24. [Online]. Available: <http://kamalmarhubi.com/blog/2015/09/06/kubernetes-from-the-ground-up-the-api-server/>
- [19] T. Almeida and J. M. G. Hidalgo, "Sms spam collection," Web Page, accessed: 2017-04-12. [Online]. Available: <http://dcomp.sor.ufscar.br/talmeida/smssпамcollection/>
- [20] Wikipedia, "Support vector machine," Web Page, accessed: 2017-05-13. [Online]. Available: [https://en.wikipedia.org/wiki/Support\\_vector\\_machine](https://en.wikipedia.org/wiki/Support_vector_machine)
- [21] "Introduction to support vector machines - opencv 2.4.13.2 documentation," Web Page, accessed: 2017-05-13. [Online]. Available: [http://docs.opencv.org/2.4/doc/tutorials/ml/introduction\\_to\\_svm/introduction\\_to\\_svm.html](http://docs.opencv.org/2.4/doc/tutorials/ml/introduction_to_svm/introduction_to_svm.html)
- [22] P.-N. Tan, M. Steinbach, and V. Kumar, *Introduction to Data Mining*, ser. Pearson International Edition. Pearson Addison Wesley, 2006. [Online]. Available: [https://books.google.com/books?id=\\_XdrQgAACAAJ](https://books.google.com/books?id=_XdrQgAACAAJ)
- [23] Wikipedia, "Polynomial kernel," Web Page, accessed: 2017-05-13. [Online]. Available: [https://en.wikipedia.org/wiki/Polynomial\\_kernel](https://en.wikipedia.org/wiki/Polynomial_kernel)
- [24] Wikipedia, "Euclidean distance - wikipedia," Web Page, accessed: 2017-06-15. [Online]. Available: [https://en.wikipedia.org/wiki/Euclidean\\_distance](https://en.wikipedia.org/wiki/Euclidean_distance)
- [25] T. Srivastava, "Introduction to knn, k-nearest neighbors : Simplified," Web Page, Oct 2014, accessed: 2017-04-18. [Online]. Available: <https://www.analyticsvidhya.com/blog/2014/10/introduction-k-neighbours-algorithm-clustering/>
- [26] Wikipedia, "Adaboost - wikipedia," Web Page, accessed: 2017-05-27. [Online]. Available: <https://en.wikipedia.org/wiki/AdaBoost>
- [27] Wikipedia, "Curse of dimensionality - wikipedia," Web Page, accessed: 2017-08-26. [Online]. Available: [https://en.wikipedia.org/wiki/Curse\\_of\\_dimensionality](https://en.wikipedia.org/wiki/Curse_of_dimensionality)
- [28] J. Provost, "Naive-Bayes vs. Rule-Learning in Classification of Email," in *Artificial Intelligence Lab*. The University Of Texas at Austin: The University Of Texas, 1999, accessed: 2017-03-10. [Online]. Available: <http://mathcs.wilkes.edu/~kapolka/cs340/provost-ai-tr-99-281.pdf>
- [29] Wikipedia, "Naive bayes spam filtering," Web Page, January 2017, accessed: 2017-03-10. [Online]. Available: [https://en.wikipedia.org/wiki/Naive\\_Bayes\\_spam\\_filtering](https://en.wikipedia.org/wiki/Naive_Bayes_spam_filtering)
- [30] "Dockerfile reference | docker documentation," Web Page, accessed: 2017-05-24. [Online]. Available: <https://docs.docker.com/engine/reference/builder/>
- [31] "Explore - docker hub," Web Page, accessed: 2017-04-25. [Online]. Available: <https://hub.docker.com/explore/>
- [32] "library/r-base - docker hub," Web Page, accessed: 2017-04-17. [Online]. Available: [https://hub.docker.com/\\_/r-base/](https://hub.docker.com/_/r-base/)
- [33] "Using kubeadm to create a cluster|kubernetes," Web Page, accessed: 2017-04-24. [Online]. Available: <https://kubernetes.io/docs/setup/independent/create-cluster-kubeadm/>
- [34] "Installing addons - kubernetes," Web Page, accessed: 2017-06-15. [Online]. Available: <https://kubernetes.io/docs/concepts/cluster-administration/addons/>
- [35] "Cmd5 module," Web Page, accessed: 2017-07-15. [Online]. Available: [https://cloudmesh.github.io/classes/lesson/prg/python\\_cmd5.html](https://cloudmesh.github.io/classes/lesson/prg/python_cmd5.html)
- [36] "What is github," Web Page, accessed: 2017-06-15. [Online]. Available: <https://guides.github.com/activities/hello-world/>

## AUTHOR BIOGRAPHIES



+**Rahul Singh** received his Bachelors in Computer Engineering from University of Mumbai, India. He is currently pursuing his Masters in Computer Science at School Of Informatics and Computing, Indiana University Bloomington.



+**Sagar Vora** received his Bachelors in Computer Sciecn from University of Mumbai, India. He is currently pursuing his Masters in Computer Science at School Of Informatics and Computing, Indiana University Bloomington. During the summer, he interned at Yahoo Inc.

## 13. WORK BREAKDOWN

Rahul Singh created the code for spam detection along with its docker image. He was responsible for doing research on kubernetes and discovering the correct steps for deploying of the application on a kubernetes cluster. He has worked on sections of the report related to the application, docker and kubernetes. Application benchmarking was done in conjunction with Sagar Vora.

Sagar Vora was responsible for automating the entire deployment process using Ansible and converting the code to CMD5 module. Ansible and CMD5 sections of the report were done by him. Application benchmarking was done in conjunction with Rahul Singh.